

## CUDA Threads

- 1) Complete and run the following `vecAdd4.cu` program, and see if it works for any  $n > T$ , e.g.  $n = 1234$  and  $T = 64$ .

```
#define n 1024      // size of vectors
#define T 240      // number of threads per block

__global__ void vecAdd(int *A, int *B, int *C) {

    int i = blockIdx.x*blockDim.x + threadIdx.x;

    if (i < n)      // allows for more threads than vector elements
        C[i] = A[i] + B[i];    // some unused
}

int main (int argc, char *argv[] ) {

    int blocks = (n + T - 1) / T;    // efficient way of rounding to next integer
    ...
    vecAdd<<<blocks, T>>>>(devA, devB, devC);
    ...

}
```

- 2) Create a CUDA program to fill in the array `A[256]` using 4 thread blocks. Each block has 64 threads. Each element `A[i] = i` as shown below.

0	1	2	3	4	5	6	.....	253	254	255
---	---	---	---	---	---	---	-------	-----	-----	-----

- 3) Create and run the `matmul2.cu` program which multiplies two square matrices. *Width* is strictly a multiple of *TILE\_WIDTH*.

```
#include <stdio.h>

#define Width 32      // size of Width x Width matrix
#define TILE_WIDTH 16

__global__ void MatrixMulKernel(float* Md, float* Nd, float* Pd, int ncols) {

    int row = blockIdx.y*blockDim.y + threadIdx.y;
    int col = blockIdx.x*blockDim.x + threadIdx.x;

    // Pvalue is used to store the element of the output matrix
    // that is computed by the thread

    float Pvalue = 0;
    for (int k = 0; k < ncols; ++k) {
        float Melement = Md[row*ncols+k];
        float Nelement = Nd[k*ncols+col];
        Pvalue += Melement * Nelement;
    }

    Pd[row*ncols+col] = Pvalue;
}
```

```

int main (int argc, char *argv[] ) {

    int i,j;
    int size = Width * Width * sizeof(float);
    float M[Width][Width], N[Width][Width], P[Width][Width];
    float* Md, *Nd, *Pd;

    for (i=0; i < Width; i++) {
        for (j=0; j < Width; j++) {
            M[i][j] = 1; N[i][j] = 2;
        }
    }

    cudaMalloc( (void**)&Md, size);
    cudaMalloc( (void**)&Nd, size);
    cudaMalloc( (void**)&Pd, size);

    cudaMemcpy( Md, M, size, cudaMemcpyHostToDevice);
    cudaMemcpy( Nd, N, size, cudaMemcpyHostToDevice);

    // Setup the execution configuration
    dim3 dimBlock(TILE_WIDTH, TILE_WIDTH);
    dim3 dimGrid(Width/TILE_WIDTH, Width/TILE_WIDTH);

    // Launch the device computation threads!
    MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);

    // Read P from the device
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);

    // Free device matrices
    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);

    for (i=0; i < Width; i++) {
        for (j=0; j < Width; j++) {
            printf("%.2f ", P[i][j]);
        }
        printf("\n");
    }
}

```

- 4) Modify the program `matmul2.cu` into `matmul3.cu` to work with square matrices of which the matrix *Width* not necessary a multiple of *TILE\_WIDTH*. For example, *Width* = 32 but *TILE\_WIDTH*=5.