

Advanced Topics in CUDA

Sudsanguan Ngamsuriyaroj
Ekasit Kijisipongse
Ittipon Rassameeroj

Semester 1/2022

Topics

- ▶ Sorting in CUDA
- ▶ CUDA reduction
- ▶ Error checking on CUDA operations
- ▶ CUDA Memory

Rank Sort

- ▶ Number of numbers that are smaller than each selected number counted. This count provides the position of selected number in sorted list; that is, its “rank.”

```
for (i = 0; i < n; i++) {      /* for each number */
    x = 0;
    for (j = 0; j < n; j++)    /* count number less than it */
        if (a[i] > a[j]) x++;
    b[x] = a[i];               /* copy number into correct place */
}
```

- ▶ Assume no duplicated numbers.
- ▶ Overall sequential sorting time complexity of $O(n^2)$.
- ▶ Not a good sequential sorting algorithm!

Parallel Rank Sort Using n Threads

- ▶ In parallel version, one thread allocated to each number. Finds final index in $O(n)$ steps.
- ▶ With all thread operating in parallel, parallel time complexity $O(n)$, given n threads.

For each thread i , do counting for $a[i]$ in parallel

$x = 0;$

for ($j = 0; j < n; j++$) $/*$ count number less than it $*/$

 if ($a[i] > a[j]$) $x++;$

$b[x] = a[i];$ $/*$ copy no into correct place $*/$

- ▶ Where i is the global ID = $blockIdx.x * blockDim.x + threadIdx.x$

CUDA Rank Sort When $n = T$ (small n)

```
__global__ void ranksort(int *dA, int *dB) {  
    int i = threadIdx.x;          // one thread block  
    int self = dA[i];  
    int x = 0;  
    for (int j=0; j < n; j++) {  
        if (self > dA[j])  
            x++;  
    }  
    dB[x] = self;  
}
```

...

```
ranksort<<<1,T>>>>(dA,dB);      // one thread block
```

CUDA Rank Sort (n/T is an integer)

```
#include<stdio.h>

#define n 64
#define T 16

__global__ void ranksort(int *dA, int *dB) {
    int i,j, x;
    i = blockIdx.x*blockDim.x + threadIdx.x;
    int self = dA[i];

    x = 0;
    for (j=0; j < n; j++) {
        if (self > dA[j])
            x++;
    }
    dB[x] = self;
}

int main() {
    int A[n], *dA, *dB;
    int i;
    int size = n * sizeof(int);
```

```

//fill the host_array randomly (no duplicated)
for(i = 0; i < n; i++)
    A[i] = n-i;

cudaMalloc( (void**)&dA,size );
cudaMalloc( (void**)&dB,size );

cudaMemcpy(dA,A,size,cudaMemcpyHostToDevice);

dim3 dimBlock(T);
dim3 dimGrid(n/T);
ranksort<<< dimGrid, dimBlock >>>(dA,dB);

cudaMemcpy(A,dB,size,cudaMemcpyDeviceToHost);

cudaFree(dA);
cudaFree(dB);


//printf the result after sorting
for(i = 0; i < n; i++) {
    printf("%d ",A[i]);
}
printf("\n");

}

```

Reduction Operation

- ▶ Reduce all of the data in an array to a single value that contains some information from the entire array.
 - ▶ Sum, maximum element, minimum element, etc.




- ▶ Useful primitive used in lots of applications

Sequential Reduction


- ▶ Start with the first element --> partial result
- ▶ Process the next element
- ▶ $O(N)$

```
int sum = data[0];  
for (i = 1; i < N; i++) {  
    sum = sum + data[i];  
}
```




Parallel Reduction

- ▶ Pair-wise reduction in steps – Tree-like structure
- ▶ $\log_2 N$ steps
- ▶ Assume that the data size is a power of 2 and the operator used in the reduction is associative, e.g. $+$, $*$



Reduction in CUDA



Simple CUDA Reduction

```
#define n 1024          // power of 2
#define T 256           // n/T must be an integer

__global__ void reduction(int *data, int stride) {
    int tid = threadIdx.x + blockIdx.x * blockDim.x;

    int idx = 2*stride*tid;
    if (idx < n) {
        data[idx] = data[idx]+data[idx+stride];
    }
}

int main () {
    int size = n *sizeof(int);
    int a[n], sum = 0;
    int *dA;


    /* Put random numbers in a[] */

    cudaMalloc( (void**)&dA, size);
    cudaMemcpy( dA, a, size, cudaMemcpyHostToDevice);
```

Simple CUDA Reduction

```
for (int s=1; s < n; s = s*2) {  
    reduction<<<n/T, T>>>(dA,s);  
}  
  
cudaMemcpy(&sum, dA, sizeof(int), cudaMemcpyDeviceToHost);  
  
cudaFree(dA);  
  
printf("%d\n",sum);  
}
```

The first element in dA contains sum.
Other elements are partial sum.



Error Checking

```
cudaError_t error;  
...  
cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);  
  
...  
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);  
error = cudaGetLastError();  
if (error != cudaSuccess) {  
    printf("CUDA Error: %s\n", cudaGetErrorString(error));  
    return 1;  
}  
...
```

Read-modify-write problem in parallel computation

- ▶ Multiple customers booking air tickets
- ▶ Each
 - ▶ Brings up a flight seat map
 - ▶ Decides on a seat
 - ▶ Update the seat map, mark the seat as taken
- ▶ A bad outcome
 - ▶ Multiple passengers ended up booking the same seat

Race Condition in Concurrent Threads

- ▶ Threads can access (read/write) shared memory.
- ▶ Consider two threads each of which is to add one to a shared data item, x . If x was initially 0, what would the value of x be after threads 1 and 2 have completed?
- ▶ Suppose that Old and New are registers, and $\text{Mem}[x] = 0$ initially



Timing Scenario #1

Time	Thread 1	Thread 2
1	(0) $\text{Old} \leftarrow \text{Mem}[x]$	
2	(1) $\text{New} \leftarrow \text{Old} + 1$	
3	(1) $\text{Mem}[x] \leftarrow \text{New}$	
4		(1) $\text{Old} \leftarrow \text{Mem}[x]$
5		(2) $\text{New} \leftarrow \text{Old} + 1$
6		(2) $\text{Mem}[x] \leftarrow \text{New}$

- ▶ Thread 1 $\text{Old} = 0$
- ▶ Thread 2 $\text{Old} = 1$
- ▶ $\text{Mem}[x] = 2$ after the sequence

Timing Scenario #2

Time	Thread 1	Thread 2
1		(0) $\text{Old} \leftarrow \text{Mem}[x]$
2		(1) $\text{New} \leftarrow \text{Old} + 1$
3		(1) $\text{Mem}[x] \leftarrow \text{New}$
4	(1) $\text{Old} \leftarrow \text{Mem}[x]$	
5	(2) $\text{New} \leftarrow \text{Old} + 1$	
6	(2) $\text{Mem}[x] \leftarrow \text{New}$	

- ▶ Thread 1 $\text{Old} = 1$
- ▶ Thread 2 $\text{Old} = 0$
- ▶ $\text{Mem}[x] = 2$ after the sequence

Timing Scenario #3

Time	Thread 1	Thread 2
1	(0) $\text{Old} \leftarrow \text{Mem}[x]$	
2	(1) $\text{New} \leftarrow \text{Old} + 1$	
3		(0) $\text{Old} \leftarrow \text{Mem}[x]$
4	(1) $\text{Mem}[x] \leftarrow \text{New}$	
5		(1) $\text{New} \leftarrow \text{Old} + 1$
6		(1) $\text{Mem}[x] \leftarrow \text{New}$

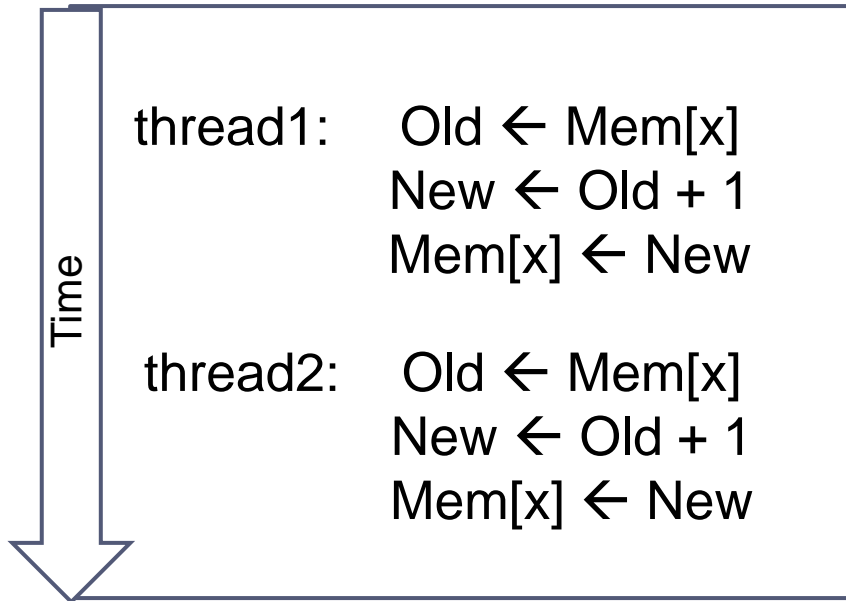
- ▶ Thread 1 $\text{Old} = 0$
- ▶ Thread 2 $\text{Old} = 0$
- ▶ $\text{Mem}[x] = 1$ after the sequence

Timing Scenario #4

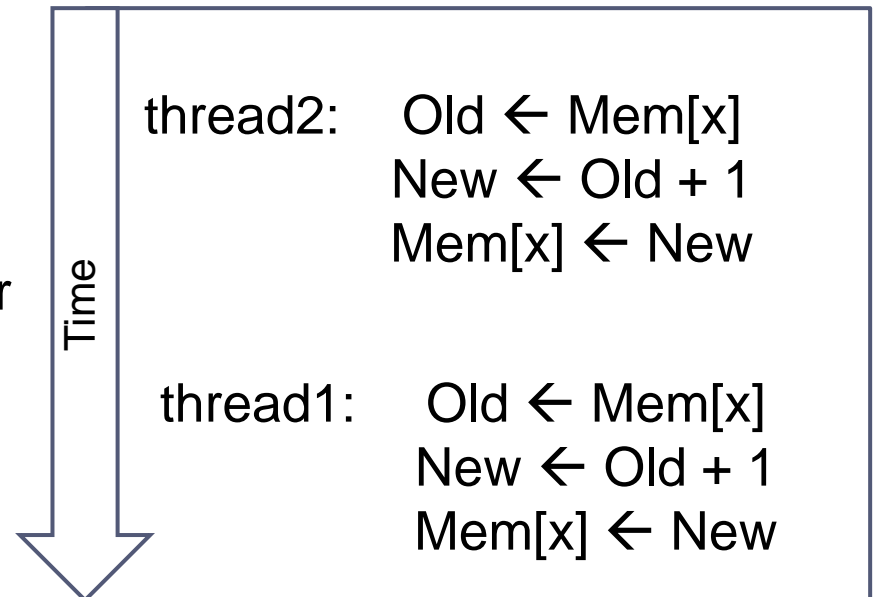
Time	Thread 1	Thread 2
1		(0) $\text{Old} \leftarrow \text{Mem}[x]$
2		(1) $\text{New} \leftarrow \text{Old} + 1$
3	(0) $\text{Old} \leftarrow \text{Mem}[x]$	
4		(1) $\text{Mem}[x] \leftarrow \text{New}$
5	(1) $\text{New} \leftarrow \text{Old} + 1$	
6	(1) $\text{Mem}[x] \leftarrow \text{New}$	

- ▶ Thread 1 $\text{Old} = 0$
- ▶ Thread 2 $\text{Old} = 0$
- ▶ $\text{Mem}[x] = 1$ after the sequence

Need Mechanism To Ensure Good Outcomes



Or



Synchronization

- ▶ Synchronization is the coordination of threads/processes to operate a system in unison
- ▶ **Mutual exclusion** is a way to ensure only one thread/process accesses a particular resource at a time.
 - ▶ **Lock** and **atomic operations** are mutual exclusion implementations
- ▶ The sections of code, called **critical sections**, must not be concurrently executed by more than one thread.
- ▶ **Barrier** synchronizes multiple threads so that any thread/process must stop at this point and cannot proceed until all other threads/processes reach this barrier.
- ▶ This concept also appears in operating systems.

Atomic Operations

- ▶ Performed by a single instruction on a memory location address
 - ▶ Read the old value, calculate a new value, and write the new value to the location
- ▶ The hardware ensures that no other threads can access the location until the atomic operation is complete
 - ▶ Any other threads that access the location will typically be held in a queue until its turn
 - ▶ All threads perform the atomic operation serially
- ▶ The advantage of atomic operations is that they are relatively quick compared to locks

AtomicAdd

- ▶ CUDA provides different types of atomic operations
 - ▶ Atomic add, sub, inc, dec, min, max, exch (exchange), CAS (compare and swap)

- Atomic Add

`atomicAdd(&x, 10)` \Rightarrow Lock x
x = x + 10
Unlock x

int atomicAdd(int **address**, int **val**);*

reads the 32-bit word **old** pointed to by **address** in global or shared memory, computes **(old + val)**, and stores the result back to memory at the same address. The function returns **old**.

- Single-precision floating-point atomic add
`float atomicAdd(float* address, float val);`

Compiler Options for Atomic Operations

- ▶ In emulation mode only, use option **sm_11** to enable atomic operations

`nvcc -deviceemu -arch=sm_11 -o progrname progrname.cu`

- ▶ In real GPU mode, atomic operations have already been enabled by default in the current NVCC compiler

A Simple AtomicAdd() Example

```
__device__ int sum = 0; // see note below
```

```
__global__ void addAll(int *data) {  
    int localVal = data[blockIdx.x * blockDim.x + threadIdx.x];  
    atomicAdd(&sum, localVal);  
}
```

Low performance since all threads sequentially add numbers to sum!!


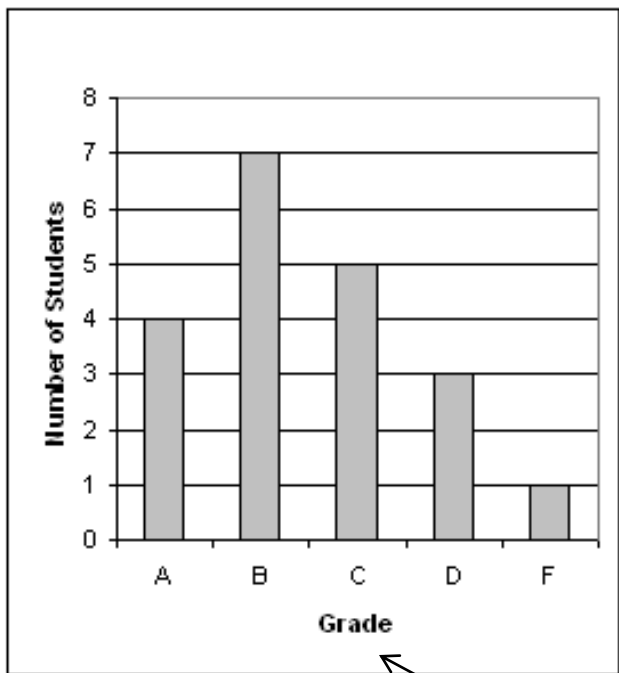
```
int main() {  
    int *dA,result;  
    ...  
    addAll<<<n/T, T>>>(dA);  
    ...  
}
```

Use reduction operation instead.

To read sum, use `cudaMemcpyFromSymbol(&result, "sum", sizeof(int), 0, cudaMemcpyDeviceToHost);`

Case Study: Histogram


- ▶ Histogram is a representation showing the frequency distribution of data, e.g. student grade or height



Bin (could be interval)

Example: Frequency of Letter

- ▶ In sentence “burberry watches” build a histogram of frequencies of each letter
 - ▶ $\text{freq}(a) = 1$, $\text{freq}(b) = 2$, $\text{freq}(c) = 1, \dots$



Histogram: Sequential Code

- For each element in the data set, use the value to identify a “bin” to increment

// assume an array of student scores from 0 to 9

```
int data[DATA_SIZE] = {2, 1, 0, 6, 2, 1, 4, 0, ...} ;
```

// Counters for 10 different scores


```
int bin[10];
```

```
for(int i = 0; i < BIN_SIZE; i++)
```

```
    bin[i] = 0;
```

```
for(int i = 0; i < DATA_SIZE; i++)
```

```
    bin[data[i]]++;
```



histogram.cu

```
#include <stdio.h>
```

```
#define n 1024
```

```
#define NUMTHREADS 256
```

```
__global__ void histogram_kernel( unsigned int *data, unsigned int *bin) {  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    if (i < n) {  
        atomicAdd( &(amp;bin[data[i]]), 1 );  
    }  
}
```

```

int main (int argc, char *argv[] ) {

    int i;
    int size = n *sizeof(int);
    unsigned int a[n];
    unsigned int bin[10];
    unsigned int *dA, *dBin;

    for (i=0; i < n; i++) {
        a[i] = i % 10;
    }
    cudaMalloc( (void**)&dA, size);
    cudaMalloc( (void**)&dBin, 10*sizeof(int));

    cudaMemcpy( dA, a, size, cudaMemcpyHostToDevice);
    cudaMemset( dBin,0, 10*sizeof(int));

    int nblocks = (n+NUMTHREADS-1)/NUMTHREADS;
    histogram_kernel<<<nblocks, NUMTHREADS>>>(dA,dBin);

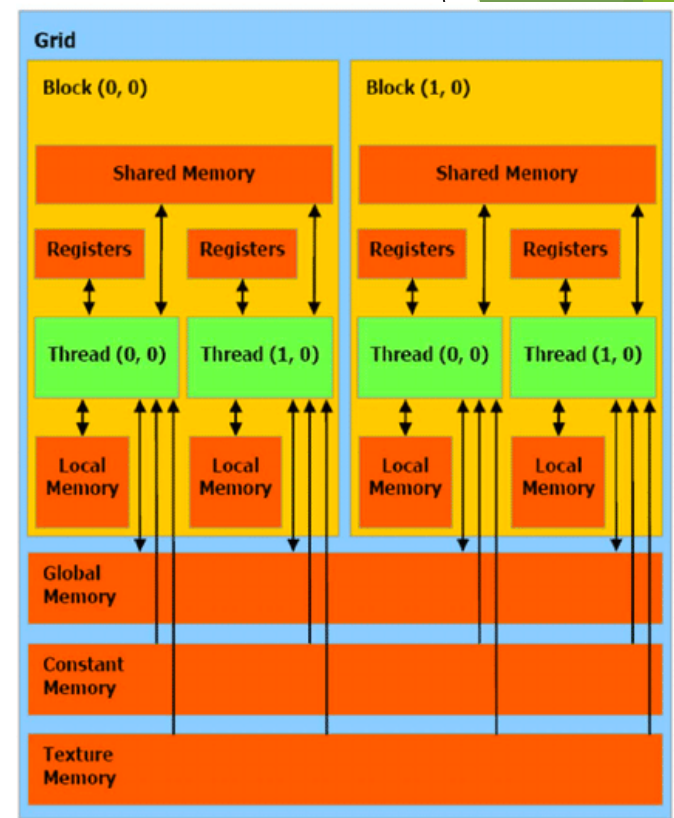
    cudaMemcpy(bin, dBin, 10*sizeof(int), cudaMemcpyDeviceToHost);
    cudaFree( dA); cudaFree( dBin);

    int count = 0;
    for (i=0; i < 10; i++) {
        printf("Freq %d = %d\n",i,bin[i]);
        count = count + bin[i];
    }
    printf("#elements = %d\n",count);
}

```

CUDA Memories


- ▶ GPUs have multiple memory spaces
- ▶ How to make the best use of the GPU memory system?
- ▶ How to deal with hardware limitation?



CUDA Memory Hierarchy

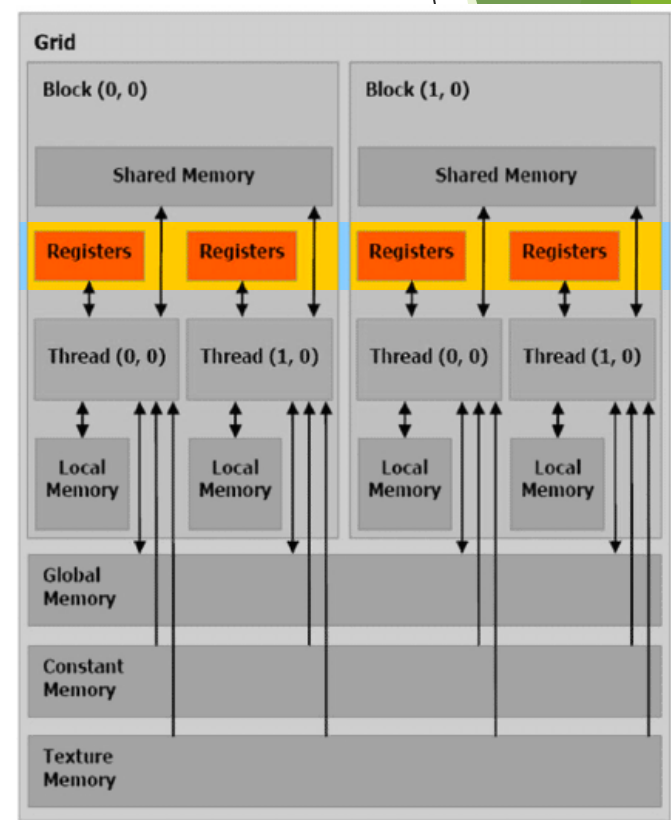
- ▶ CUDA threads may access data from multiple memory spaces
 - ▶ On-chip memories => inside SM
 - ▶ Off-chip memories => device memory

Access Level	Memory Type	Location
Thread-private	register	on-chip
	local Memory	off-chip
Thread block	shared memory	on-chip
Grid	global memory	off-chip
	constant memory	
	texture memory	



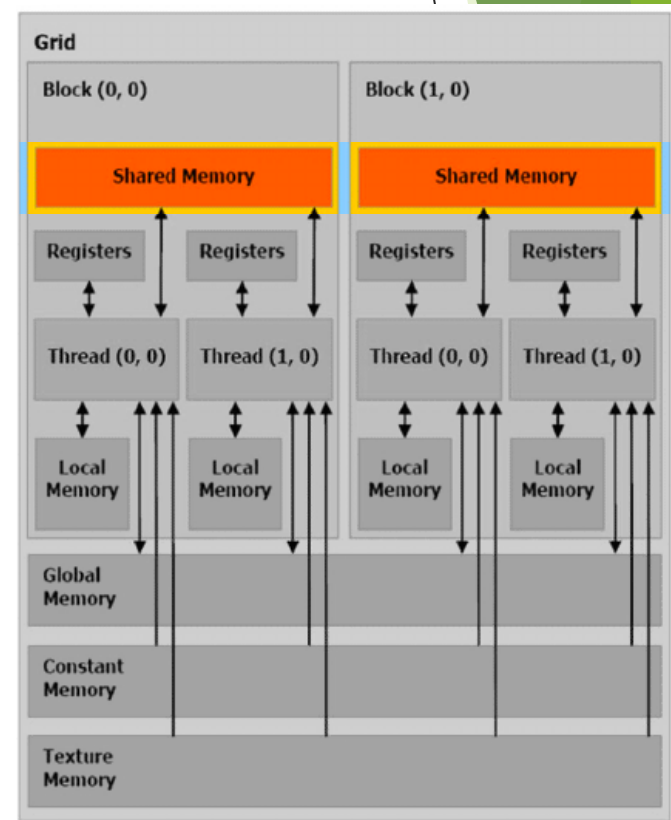
Registers

- ▶ **Fastest** memory
- ▶ Only accessible by **a thread**
- ▶ Lifetime of a thread
- ▶ Automatically allocated
 - ▶ Typically for scalar variables
- ▶ Number of registers is **very limited**
- ▶ Resides in register files
 - ▶ Shared across all active threads in an SM
- ▶ Cannot access by host



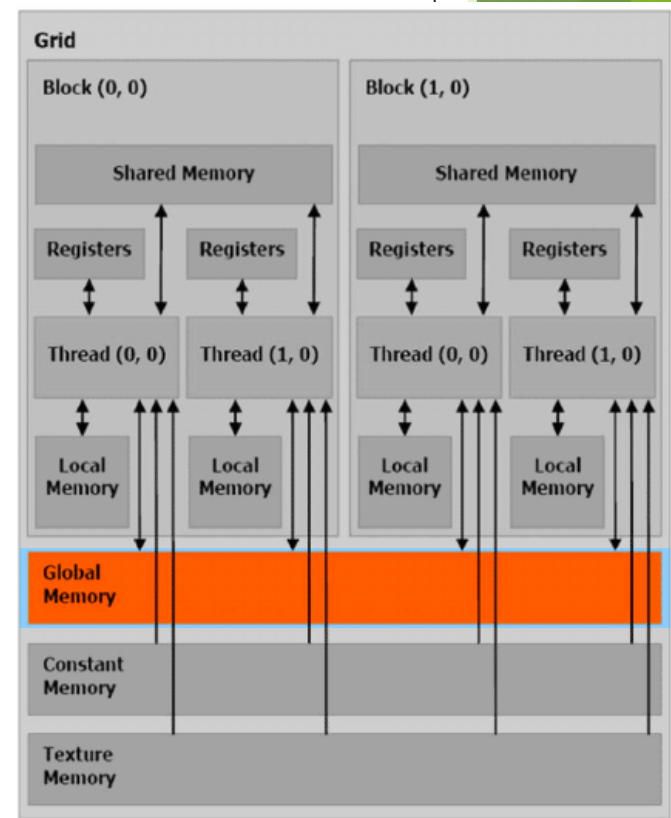
Shared Memory

- ▶ Extremely fast
- ▶ Shared across threads in the **same thread block**
- ▶ Lifetime of a kernel
- ▶ User-managed
 - ▶ Programmers must explicitly allocate shared memory in the kernel
 - ▶ `__shared__` specifier
 - ▶ Will discuss later
- ▶ Cannot access by host



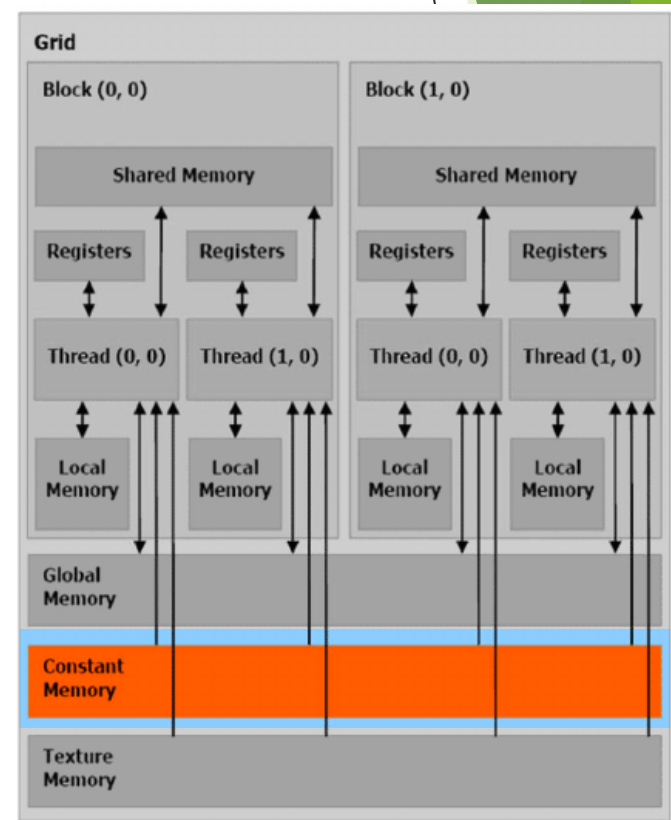
Global Memory

- ▶ Also called “**device memory**”
- ▶ Accessible by **all threads**
- ▶ Lifetime of a program
- ▶ Very high access latency
 - ▶ 400-800 clock cycles
- ▶ Potential of traffic congestion
- ▶ Can be allocated and modified by host



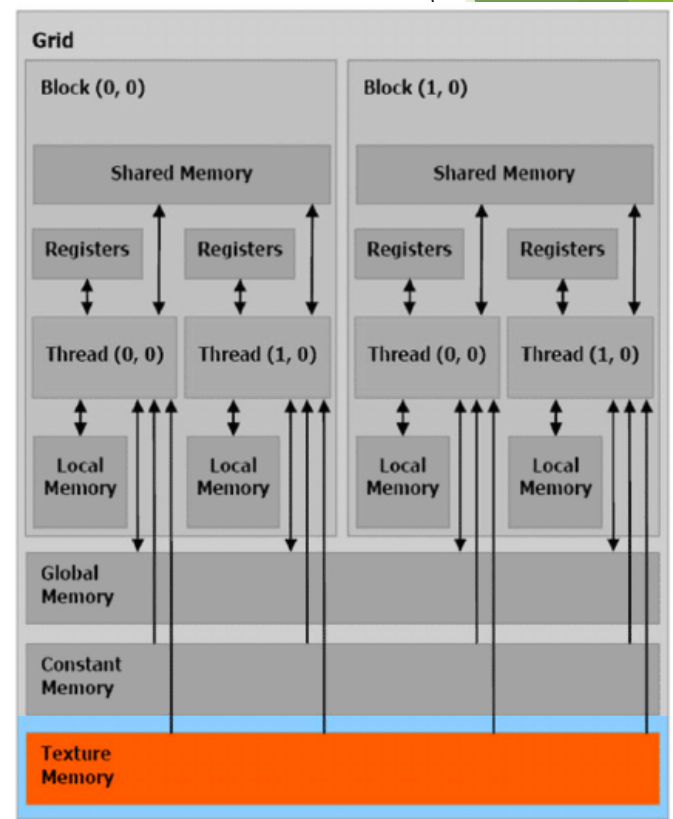

Constant Memory

- ▶ Special type of global memory
- ▶ Read-only
- ▶ Cached
- ▶ Short latency and high bandwidth when **all threads access the same location**



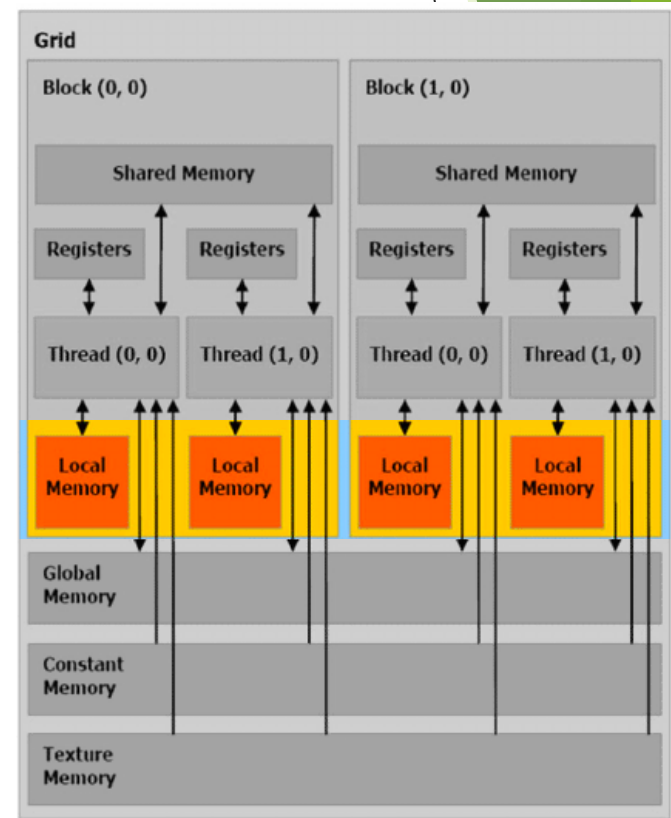
Texture Memory

- ▶ Special type of global memory
- ▶ **Read-only**
- ▶ Cached
- ▶ Optimized for 1D, 2D, or 3D spatial locality in graphic programs
 - ▶ Threads read memory is nearby location



Local Memory

- ▶ Not an actual memory space
- ▶ Similar to register
 - ▶ Only accessible by **a thread**
 - ▶ Lifetime of a thread
 - ▶ Automatically allocated
- ▶ Resides in **global memory**
- ▶ Used for
 - ▶ Storing arrays declared inside a kernel
 - ▶ Register spilling



Reducing Global Memory Traffic

- ▶ Global memory is slow
 - ▶ Can be **performance bottleneck**.
- ▶ Reducing global memory access enhances performance
- ▶ Using the memory space with caching could improve performance
 - ▶ Constant value => constant memory
 - ▶ 2D/3D memory accesses => texture memory
- ▶ What else?

Reducing Global Memory Traffic using Shared Memory

- ▶ Extremely fast, user-managed, on-chip memory.
 - ▶ ~100x faster than global memory
- ▶ Could be used as user-managed cache
 - ▶ Reduce global memory traffic
- ▶ Shared across all threads in a thread block.
 - ▶ Data could be loaded from global memory **once** and shares across all threads in the same block
 - ▶ Data will stay in the shared memory as long as
 - ▶ The thread block ends its executions
 - ▶ Programmer explicitly replace it
- ▶ Declare using `__shared__` specifier

Back to Matrix Multiplication

- ▶ **Strategy:** partition the data into subsets called “tiles”, such that each tile fits into the shared memory.
 - ▶ Data loaded inside a tile could be reused
 - ▶ Make threads that use common elements collaborate
- ▶ Each thread can load a submatrix into the shared memory before calculations
- ▶ The submatrix will be used by the thread that loaded them and other threads that share them

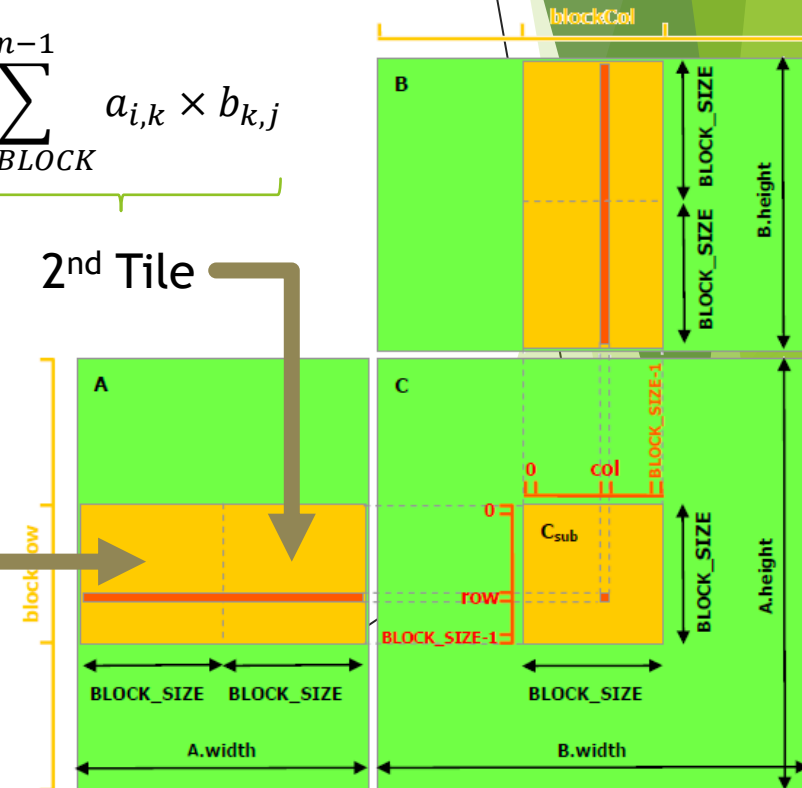
Matrix Multiplication with Shared Memory

- Looking at matrix multiplication equation

$$c_{i,j} = \sum_{k=0}^{n-1} a_{i,k} \times b_{k,j} = \underbrace{\sum_{k=0}^{BLOCK-1} a_{i,k} \times b_{k,j}}_{1^{st} \text{ Tile}} + \underbrace{\sum_{k=BLOCK}^{n-1} a_{i,k} \times b_{k,j}}_{2^{nd} \text{ Tile}}$$

1st Tile

2nd Tile



Matrix Multiplication with Shared Memory

```
__global__ void matrixMul(float* A, float* B, float* C, int width)
{
    __shared__ float As[TILE_WIDTH][TILE_WIDTH];
    __shared__ float Bs[TILE_WIDTH][TILE_WIDTH];

    int row = blockIdx.y * TILE_WIDTH + threadIdx.y;
    int col = blockIdx.x * TILE_WIDTH + threadIdx.x;

    float c_val = 0.0f;

    for(int i = 0; i < width/TILE_WIDTH; i++){
        As[threadIdx.y][threadIdx.x] = A[row * width + (i * TILE_WIDTH + threadIdx.x)];
        Bs[threadIdx.y][threadIdx.x] = B[(i * TILE_WIDTH + threadIdx.y) * width + col];
        __syncthreads();

        for(int k = 0; k < TILE_WIDTH; k++){
            c_val += As[threadIdx.y][k] * Bs[k][threadIdx.x];
            __syncthreads();
        }
        C[row * width + col] = c_val;
    }
}
```

Matrix Multiplication with Shared Memory

```
__global__ void matrixMul(float* A, float* B, float* C, int width)
{
    __shared__ float As[TILE_WIDTH][TILE_WIDTH];
    __shared__ float Bs[TILE_WIDTH][TILE_WIDTH];

    int row = blockIdx.y * TILE_WIDTH + threadIdx.y;
    int col = blockIdx.x * TILE_WIDTH + threadIdx.x;

    float c_val = 0.0f;

    for(int i = 0; i < width/TILE_WIDTH; i++){
        As[threadIdx.y][threadIdx.x] = A[row * width + (i * TILE_WIDTH + threadIdx.x)];
        Bs[threadIdx.y][threadIdx.x] = B[(i * TILE_WIDTH + threadIdx.y) * width + col];
        __syncthreads();

        for(int k = 0; k < TILE_WIDTH; k++){
            c_val += As[threadIdx.y][k] * Bs[k][threadIdx.x];
            __syncthreads();
        }
        C[row * width + col] = c_val;
    }
}
```

Allocate As and Bs in shared memory

Matrix Multiplication with Shared Memory

```
__global__ void matrixMul(float* A, float* B, float* C, int width)
{
    __shared__ float As[TILE_WIDTH][TILE_WIDTH];
    __shared__ float Bs[TILE_WIDTH][TILE_WIDTH];

    int row = blockIdx.y * TILE_WIDTH + threadIdx.y;
    int col = blockIdx.x * TILE_WIDTH + threadIdx.x;

    float c_val = 0.0f;

    for(int i = 0; i < width/TILE_WIDTH; i++){
        As[threadIdx.y][threadIdx.x] = A[row * width + (i * TILE_WIDTH + threadIdx.x)];
        Bs[threadIdx.y][threadIdx.x] = B[(i * TILE_WIDTH + threadIdx.y) * width + col];
        __syncthreads();

        for(int k = 0; k < TILE_WIDTH; k++)
            c_val += As[threadIdx.y][k] * Bs[k][threadIdx.x];
        __syncthreads();
    }
    C[row * width + col] = c_val;
}
```

Iterate through each tile

Matrix Multiplication with Shared Memory

```
__global__ void matrixMul(float* A, float* B, float* C, int width)
{
    __shared__ float As[TILE_WIDTH][TILE_WIDTH];
    __shared__ float Bs[TILE_WIDTH][TILE_WIDTH];

    int row = blockIdx.y * TILE_WIDTH + threadIdx.y;
    int col = blockIdx.x * TILE_WIDTH + threadIdx.x;

    float c_val = 0.0f;

    for(int i = 0; i < width/TILE_WIDTH; i++){
        As[threadIdx.y][threadIdx.x] = A[row * width + (i * TILE_WIDTH + threadIdx.x)];
        Bs[threadIdx.y][threadIdx.x] = B[(i * TILE_WIDTH + threadIdx.y) * width + col];
        __syncthreads();
        for(int k = 0; k < TILE_WIDTH; k++){
            c_val += As[threadIdx.y][k] * Bs[k][threadIdx.x];
            __syncthreads();
        }
        C[row * width + col] = c_val;
    }
}
```

Make sure that all
elements are loaded

Matrix Multiplication with Shared Memory

```
__global__ void matrixMul(float* A, float* B, float* C, int width)
{
    __shared__ float As[TILE_WIDTH][TILE_WIDTH];
    __shared__ float Bs[TILE_WIDTH][TILE_WIDTH];

    int row = blockIdx.y * TILE_WIDTH + threadIdx.y;
    int col = blockIdx.x * TILE_WIDTH + threadIdx.x;

    float c_val = 0.0f;

    for(int i = 0; i < width/TILE_WIDTH; i++){
        As[threadIdx.y][threadIdx.x] = A[row * width + (i * TILE_WIDTH + threadIdx.x)];
        Bs[threadIdx.y][threadIdx.x] = B[(i * TILE_WIDTH + threadIdx.y) * width + col];
        __syncthreads();

        for(int k = 0; k < TILE_WIDTH; k++){
            c_val += As[threadIdx.y][k] * Bs[k][threadIdx.x];
            __syncthreads();
        }
        C[row * width + col] = c_val;
    }
}
```

Make sure calculation
are completed before
loading new values⁴⁸

Matrix Multiplication with Shared Memory

- ▶ GPU on-chip resources are limited
- ▶ Each SMs has limited shared memory space
- ▶ Shared memory space is also shared by thread blocks
- ▶ What if a thread block allocate large shared memory ?
 - ▶ Maxwell SM can execute 32 thread blocks maximum
 - ▶ Maxwell SM has 64KB shared memory space
 - ▶ If a thread block requires 8KB of shared memory
 - ▶ E.g. `TILE_WIDTH` is 32
 - ▶ $32 \times 32 \times 4 = 4K$ for As and Bs
 - ▶ Only 8 thread blocks can reside in an SM at a time

Q & A