

Hadoop



Sudsanguan Ngamsuriyaroj
Ekasit Kijisipongse
Ittipon Rassameeroj

Semester 1/2022

Topics

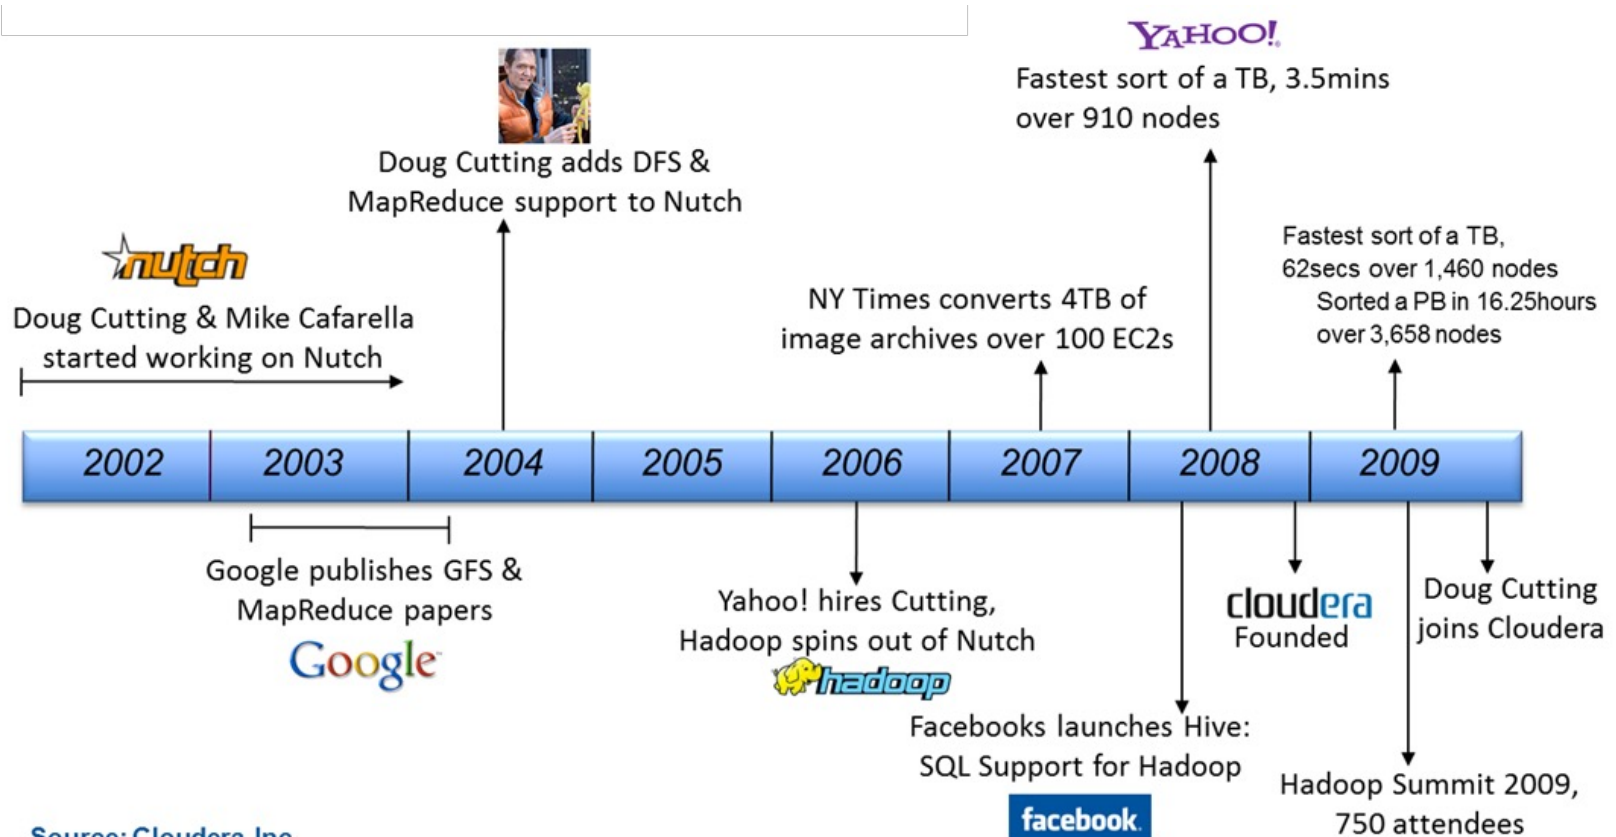
- What is Hadoop
- Hadoop Distributed File System
- Hadoop MapReduce
- Programming in MapReduce
- Hadoop Ecosystem

What is ***hadoop***

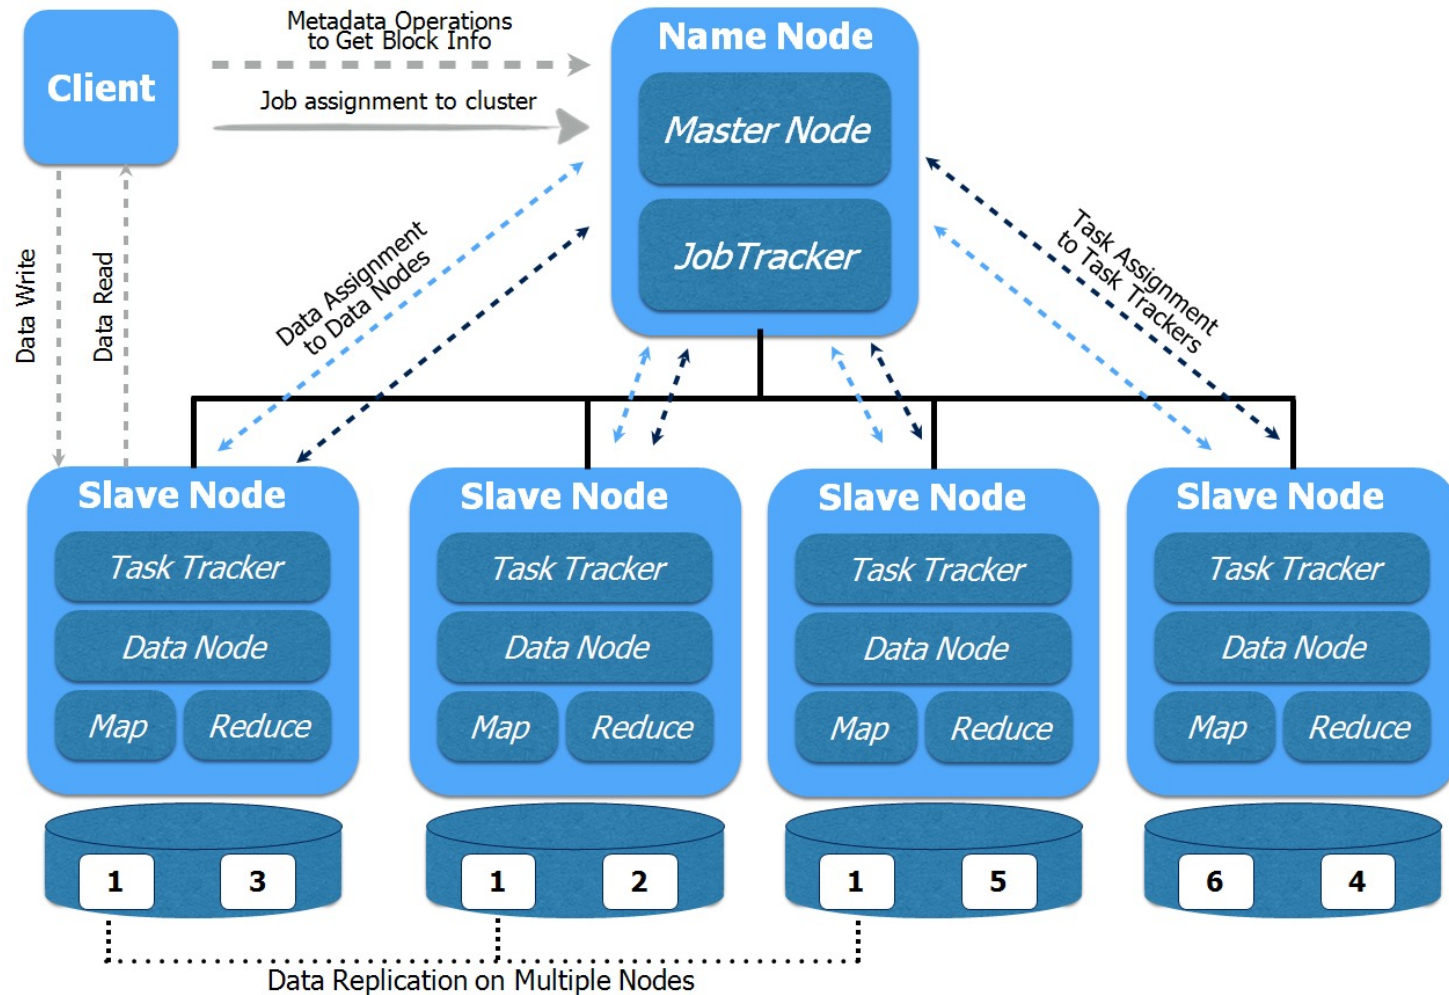
- How to process Big data with reasonable cost and time?
- Hadoop is an open-source software that supports Big data applications, licensed under the Apache v2 license
- Design Principles
 - Facilitate the storage and processing of large and/or rapidly growing data sets, both structured and non-structured data
 - Simple programming models
 - Optimized for large and very large data
 - Highly scalable
 - Use commodity (cheap!) hardware
 - Fault-tolerance
 - Move computation rather than data

Hadoop History

- Starting by Cutting at Yahoo based on Google's publication in 2003 and 2004
- Developers works at several organizations, e.g facebook, Cloudera, Yahoo



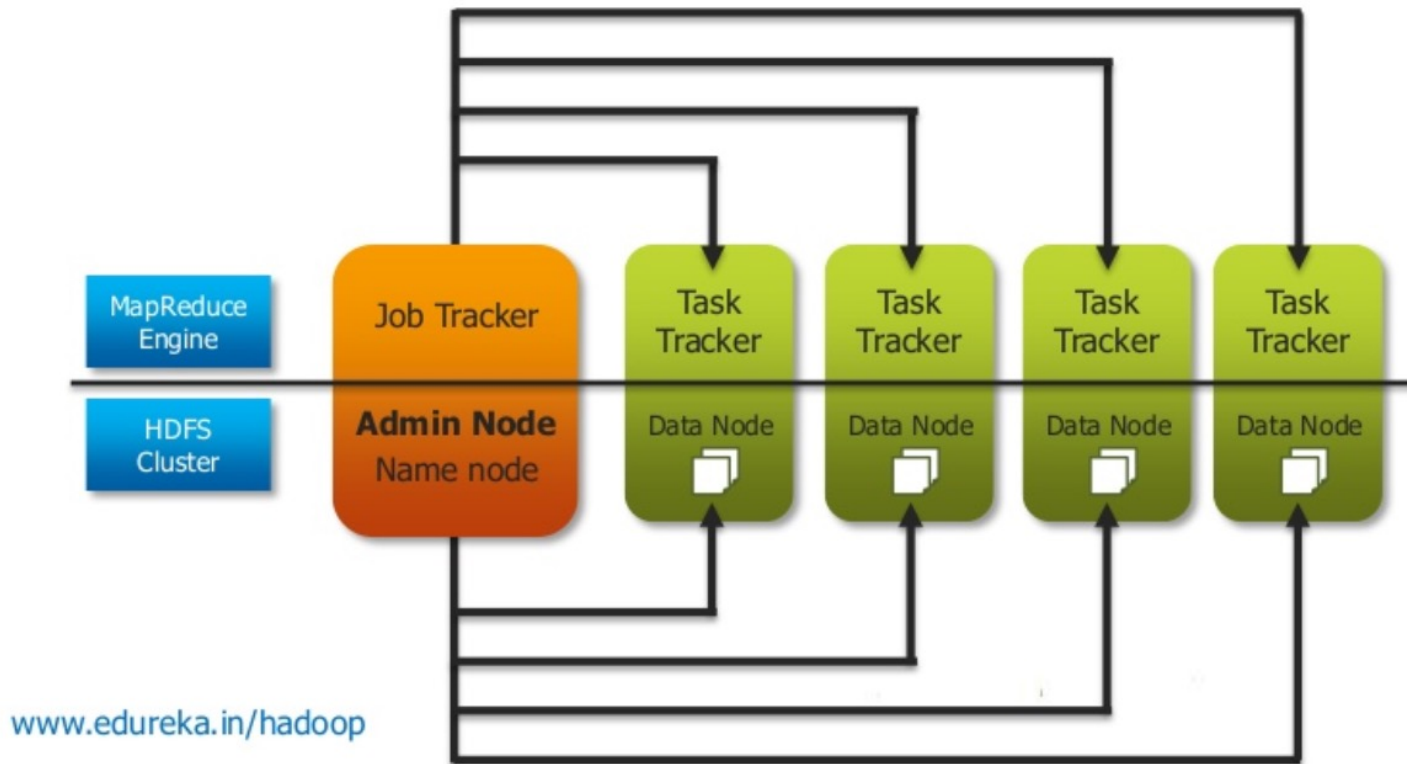
Batch Processing in Hadoop



Hadoop Core Architecture

Store Process

Hadoop = HDFS + MapReduce




Hadoop Distributed File System (HDFS)

- Very Large Distributed File System (not Databases)
 - 10K nodes, 100 million files, 10 PB
- Assumes Commodity Hardware
 - Files are replicated to handle hardware failure
 - Detect failures and recovers from them
- Optimized for Large Streaming Reads of Files
 - Provides very high aggregate bandwidth
- Files are broken up into blocks
 - Typically 64MB block size
- Each block replicated on multiple DataNodes
- Not support random write! (only created and appended)

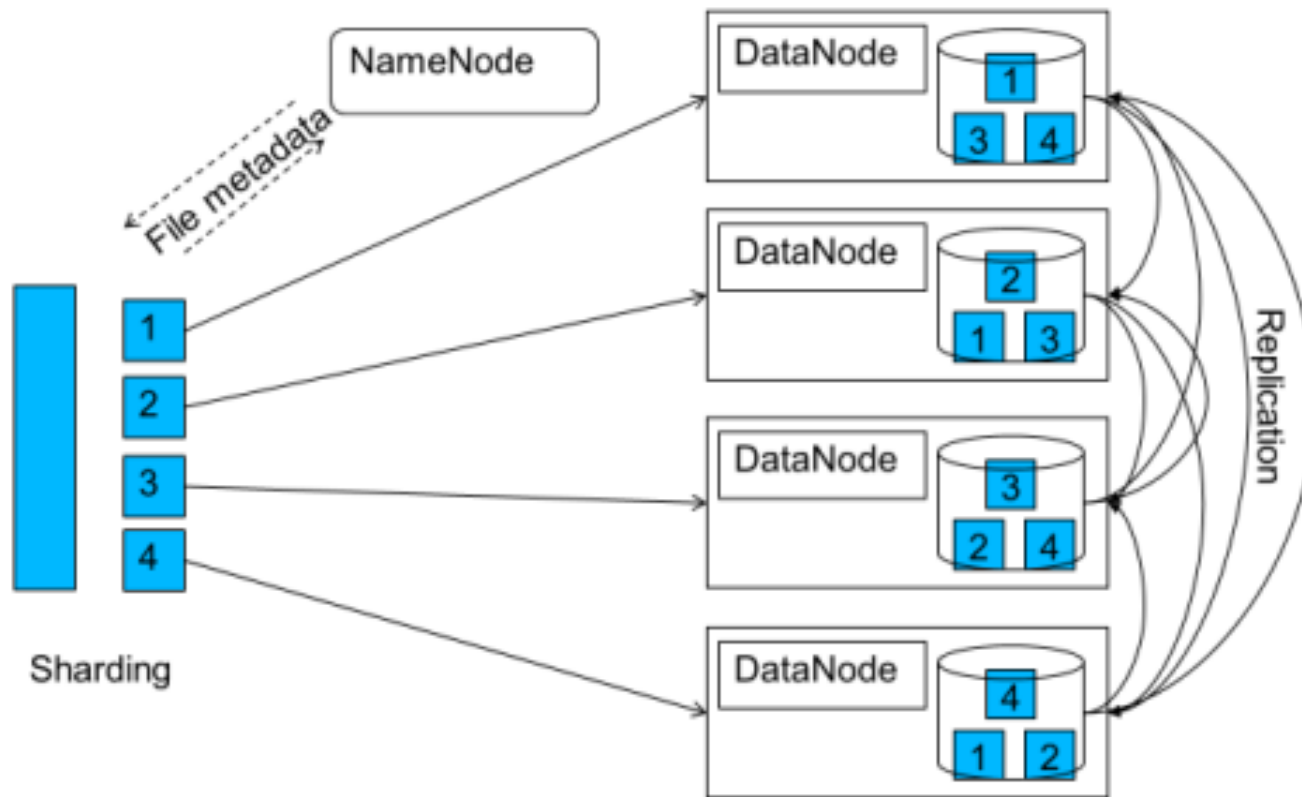
How Files are Stored

- Files are split into blocks and distributed to datanodes
- Each block is replicated to multiple nodes (default to 3)



HDFS Architecture

- NameNode + DataNode



Namenode

- Maintains the HDFS namespace, filesystem tree and metadata
- Maintains the mapping from each file to the list of blockIDs where the file is
- Metadata mapping is maintained in memory as well as persisted on disk
- Maintains in memory the locations of each block. (Block to datanode mapping)
- Memory requirement: ~150 bytes/file
- Issues instructions to datanode to create/replicate/delete blocks

Datanodes

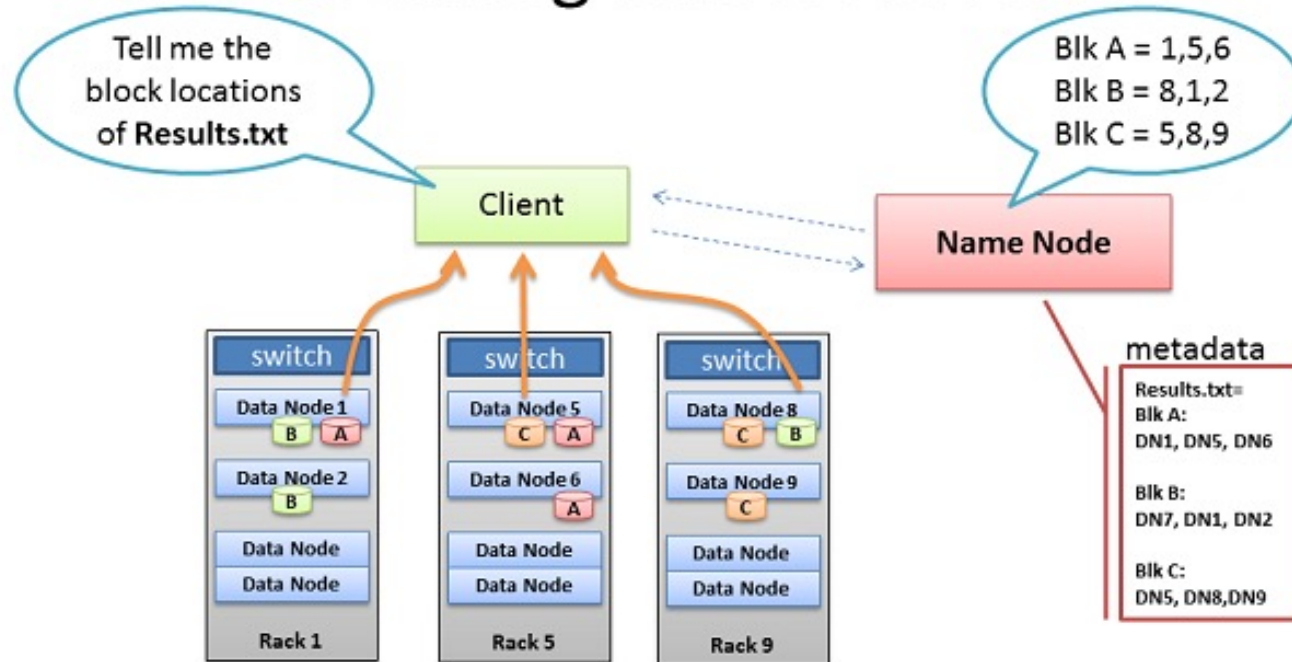
- Serve as storage for data blocks
- No metadata
- Report all blocks to namenode at startup (BlockReport)
- Sends periodic "heartbeat" to Namenode
- Serves read, write requests, performs block creation, deletion, and replication upon instruction from Namenode
- User data never flows through the NameNode

Replication and Rack-awareness

- Replication in Hadoop is at the block level
- Each block of data will be replicated to multiple machines to prevent the failure of one machine from losing all copies of data
- Unfortunate if all copies of data happened to be located on machines in the same rack, and that rack experiences a failure?
- Replication is "Rack-aware"
- Reading and writing on HDFS also makes use of rack-awareness

Reading Files from HDFS

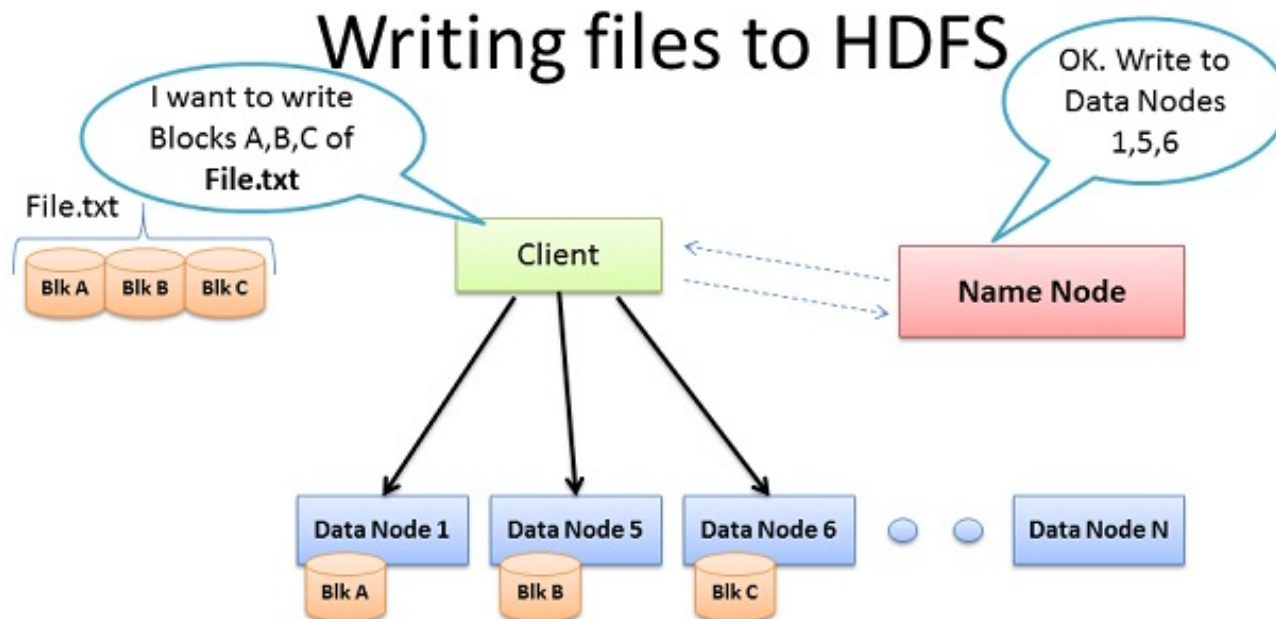
Client reading files from HDFS



- Client receives Data Node list for each block
- Client picks first Data Node for each block
- Client reads blocks sequentially

BRAD HEDLUND .com

Writing Files to HDFS

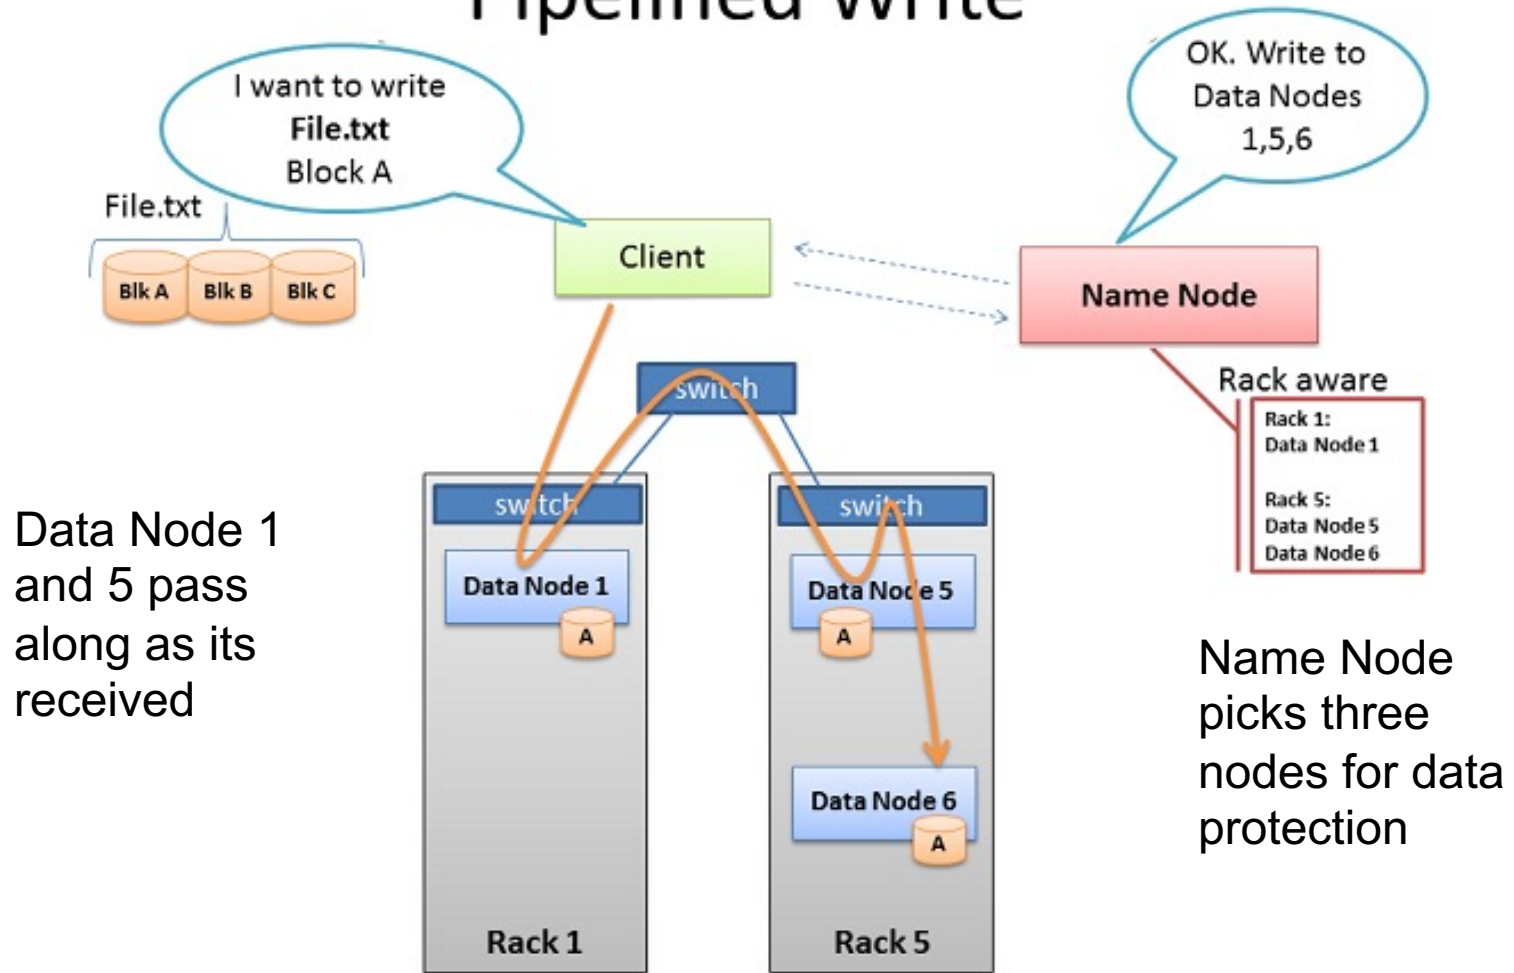


- Client consults Name Node
- Client writes block directly to one Data Node
- Data Nodes replicates block
- Cycle repeats for next block

BRAD HEDLUND .com

Pipelined Write

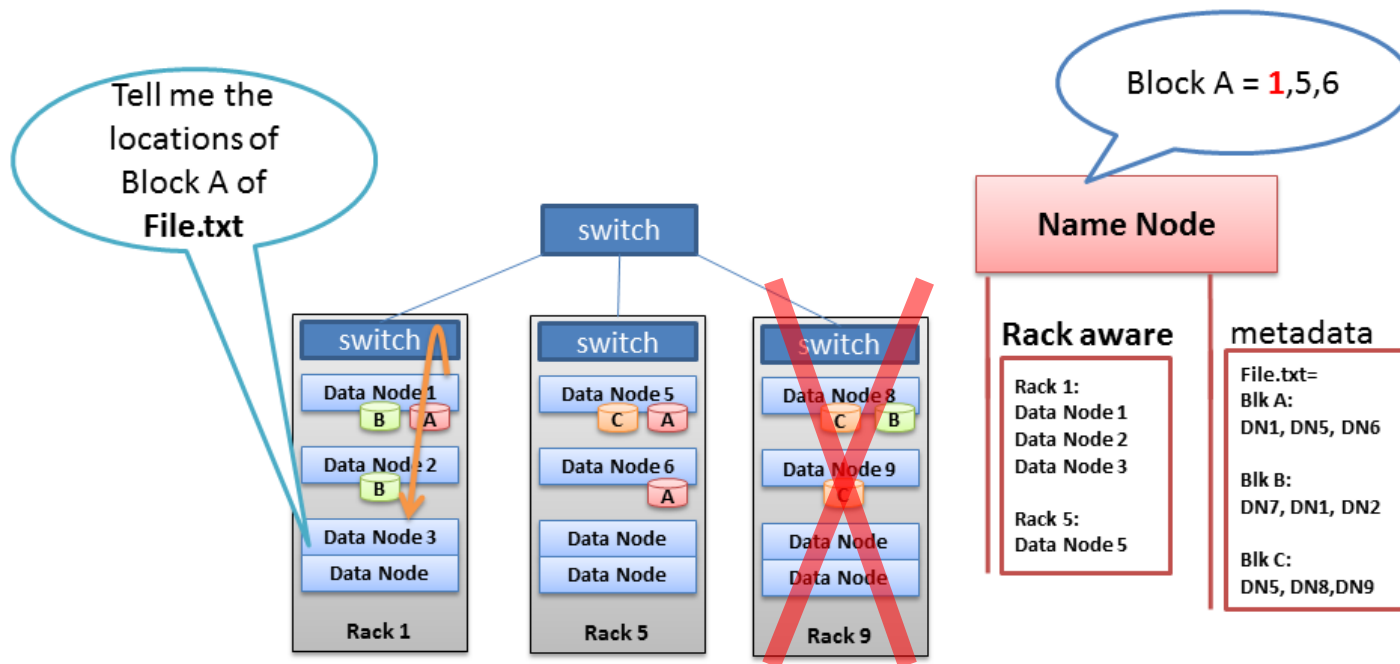
Pipelined Write



BRAD HEDLUND .com

Fault-Tolerance

- Never lose data even if an entire rack fails

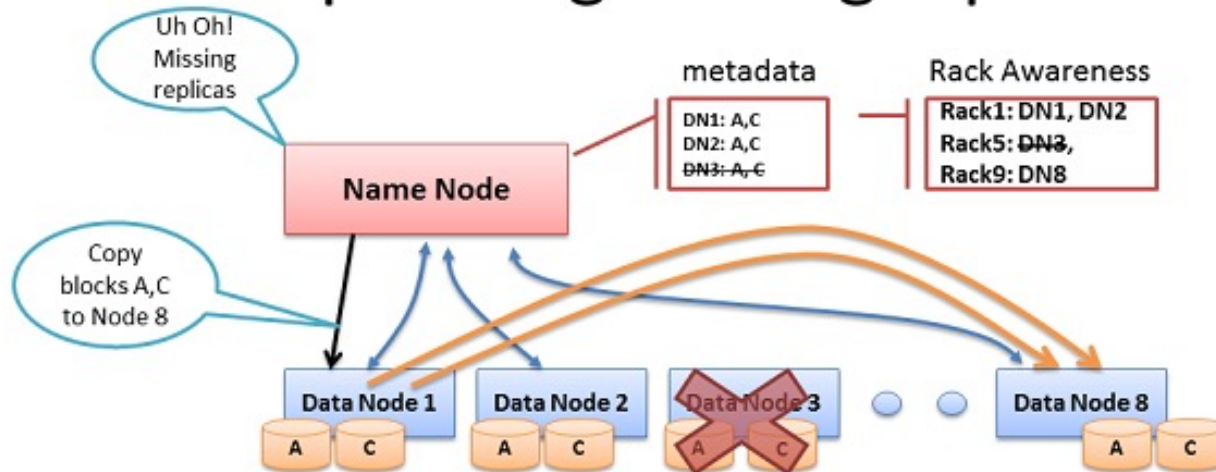


- Name Node provides rack local Nodes first
- Leverage in-rack bandwidth, single hop

BRAD HEDLUND .com

Recover from Fault

Re-replicating missing replicas



- Missing Heartbeats signify lost Nodes
- Name Node consults metadata, finds affected data
- Name Node consults Rack Awareness script
- Name Node tells a Data Node to re-replicate


BRAD HEDLUND .com

MapReduce

- Pioneered by Google to process 20 petabytes of data per day
- Simple parallel programming model designed for processing **large datasets** on a large cluster of machines
 - scalability and fault-tolerance
- Programmer specifies two functions: **Map** and **Reduce** functions.
- Input & Output: each a set of **key/value** pairs
- Many real world tasks are expressible in this model
- MPI is designed for numerical applications

MapReduce Architecture

- **Master** node runs **JobTracker**, which accepts Job requests from clients and sends tasks to **TaskTrackers**
- TaskTracker runs on **slave** nodes
- TaskTracker forks separate process for task instances



How Jobs are Executed

■ The Mapper

- Each Map task (typically) operates on a single HDFS block
- Map tasks(usually) run on the node where the block is stored



■ Shuffle and Sort

- Intermediate data from all mappers are split
- Then, consolidate and sort intermediate data to reducers
- Happens after all Map tasks are complete and before Reduce tasks start

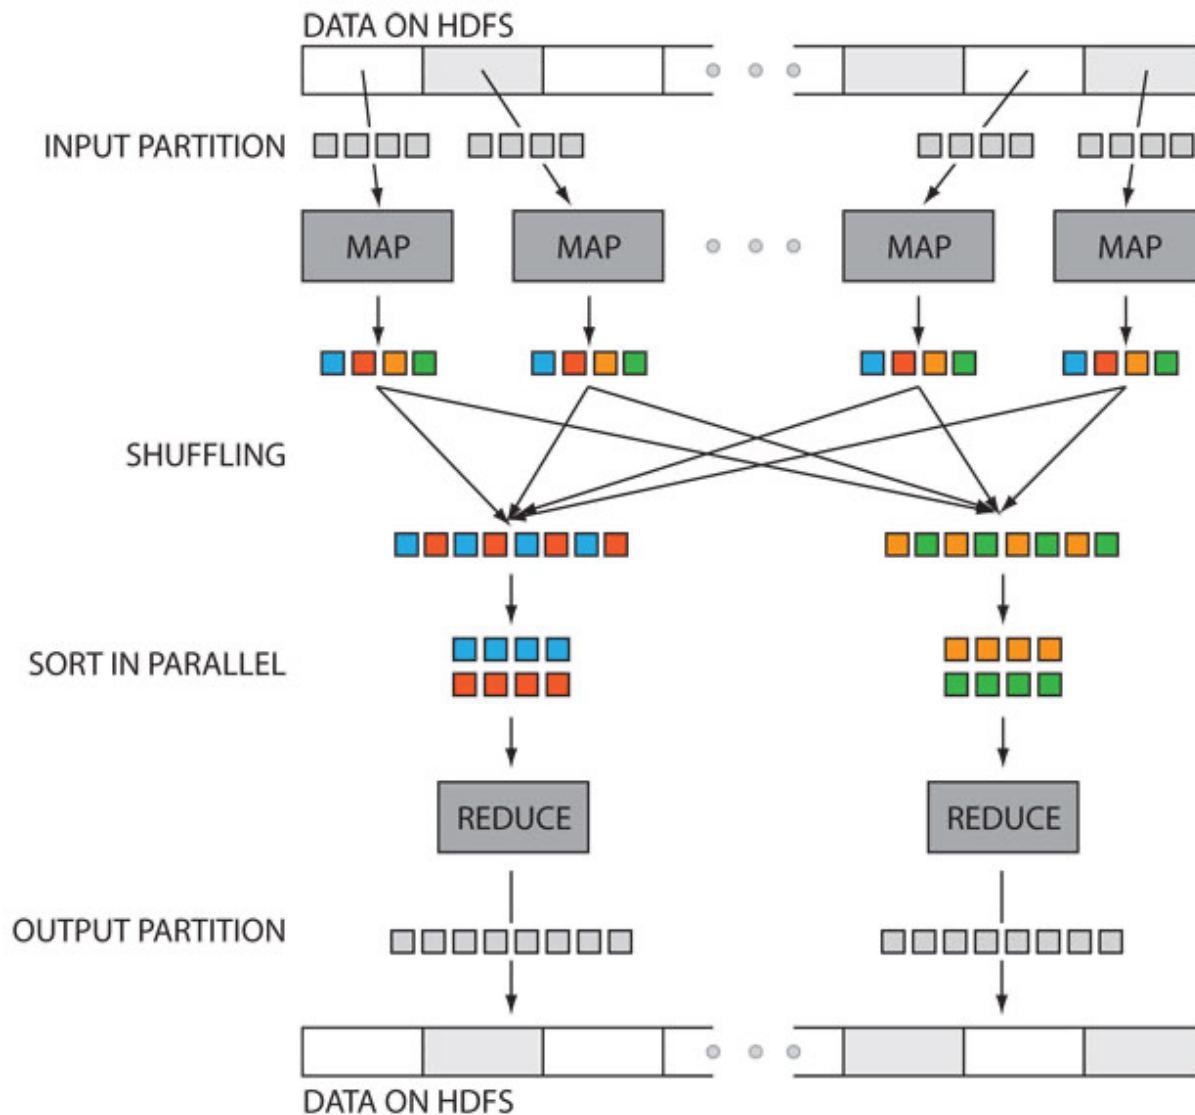


■ The Reducer

- Operates on shuffled/sorted intermediate data (Map task output)
- Produces final output



How Jobs are Executed



MapReduce Input/Output

- Input is a set of key/value (k_1, v_1) pairs
- Map converts each input to intermediate data
 $\text{map}(k_1, v_1) \rightarrow \text{list}(k_2, v_2)$
- Intermediate data are a list of (k_2, v_2) pairs
- Then, intermediate data are grouped by key
- Reduce collects and converts intermediate data into output
 $\text{reduce}(k_2, \text{list}(v_2)) \rightarrow (k_3, v_3)$
- Output is the list of (k_3, v_3) pairs
- Key cannot be null, but value can
- Example: Count word frequency in text files

Wordcount Example

Input


Line1 Hadoop uses MapReduce
Line2 There is a Map phrase
Line3
Line4 There is a Reduce phrase

Output


a 2
hadoop 1
is 2
map 1
mapreduce 1
phrase 2
reduce 1
there 2
uses 1

(This example ignores case sensitive)

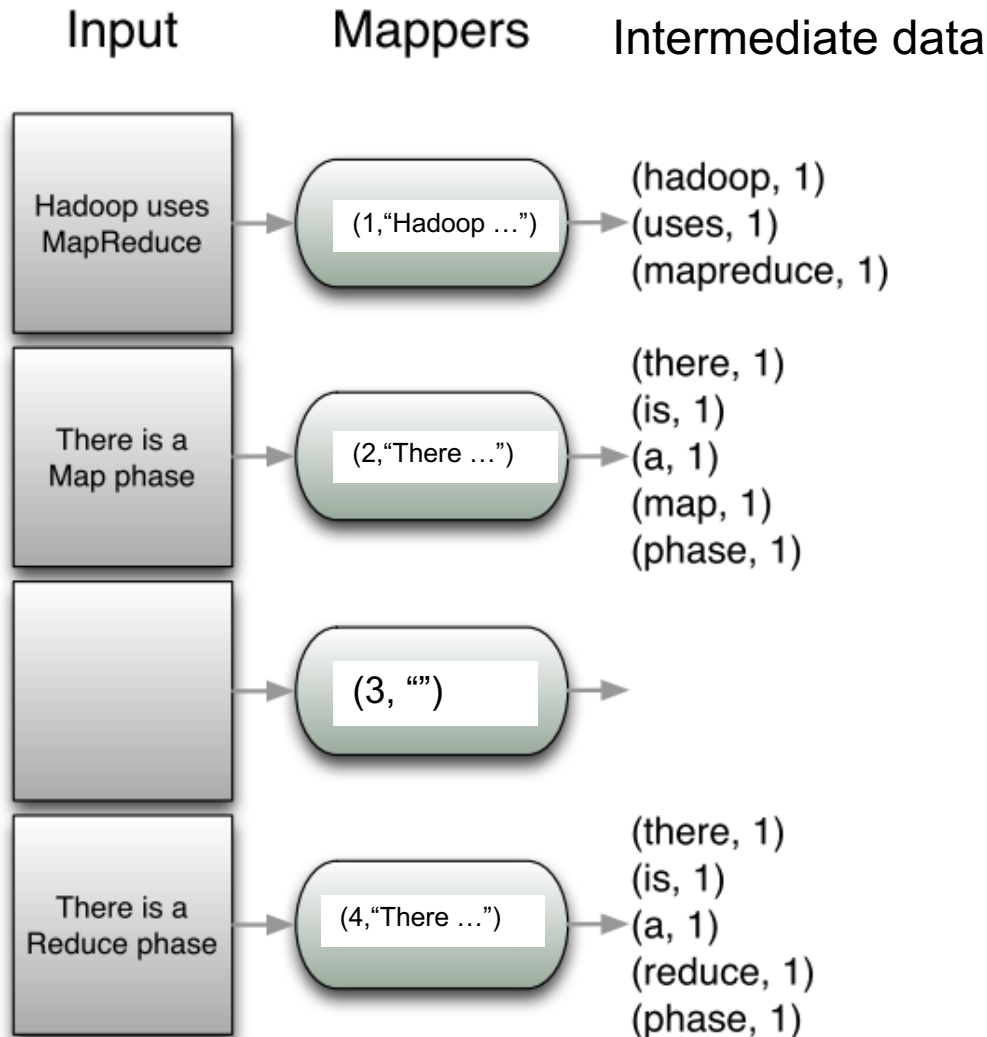
Wordcount Example



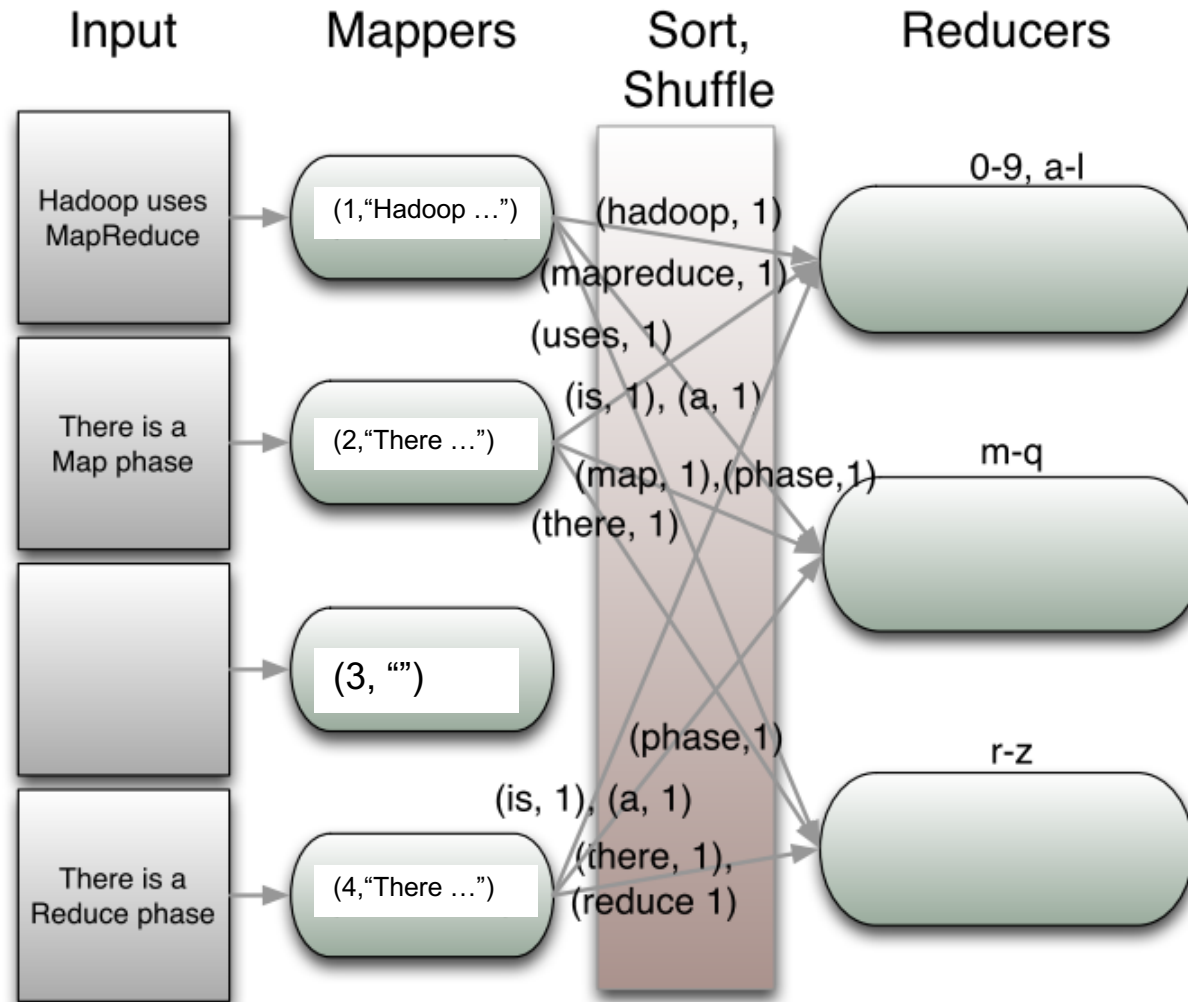
Wordcount Example



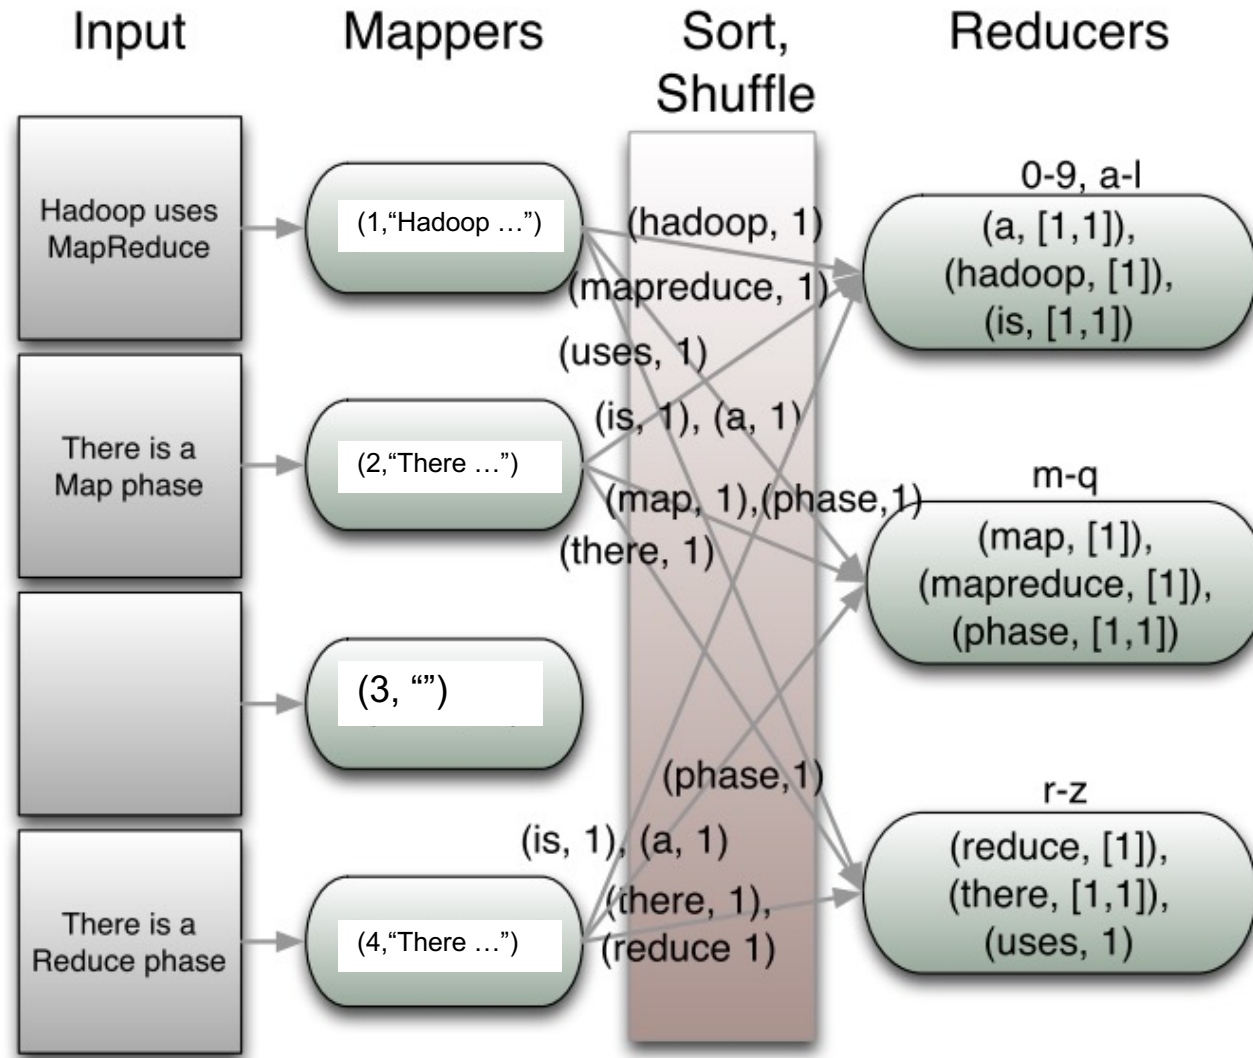
Wordcount Example



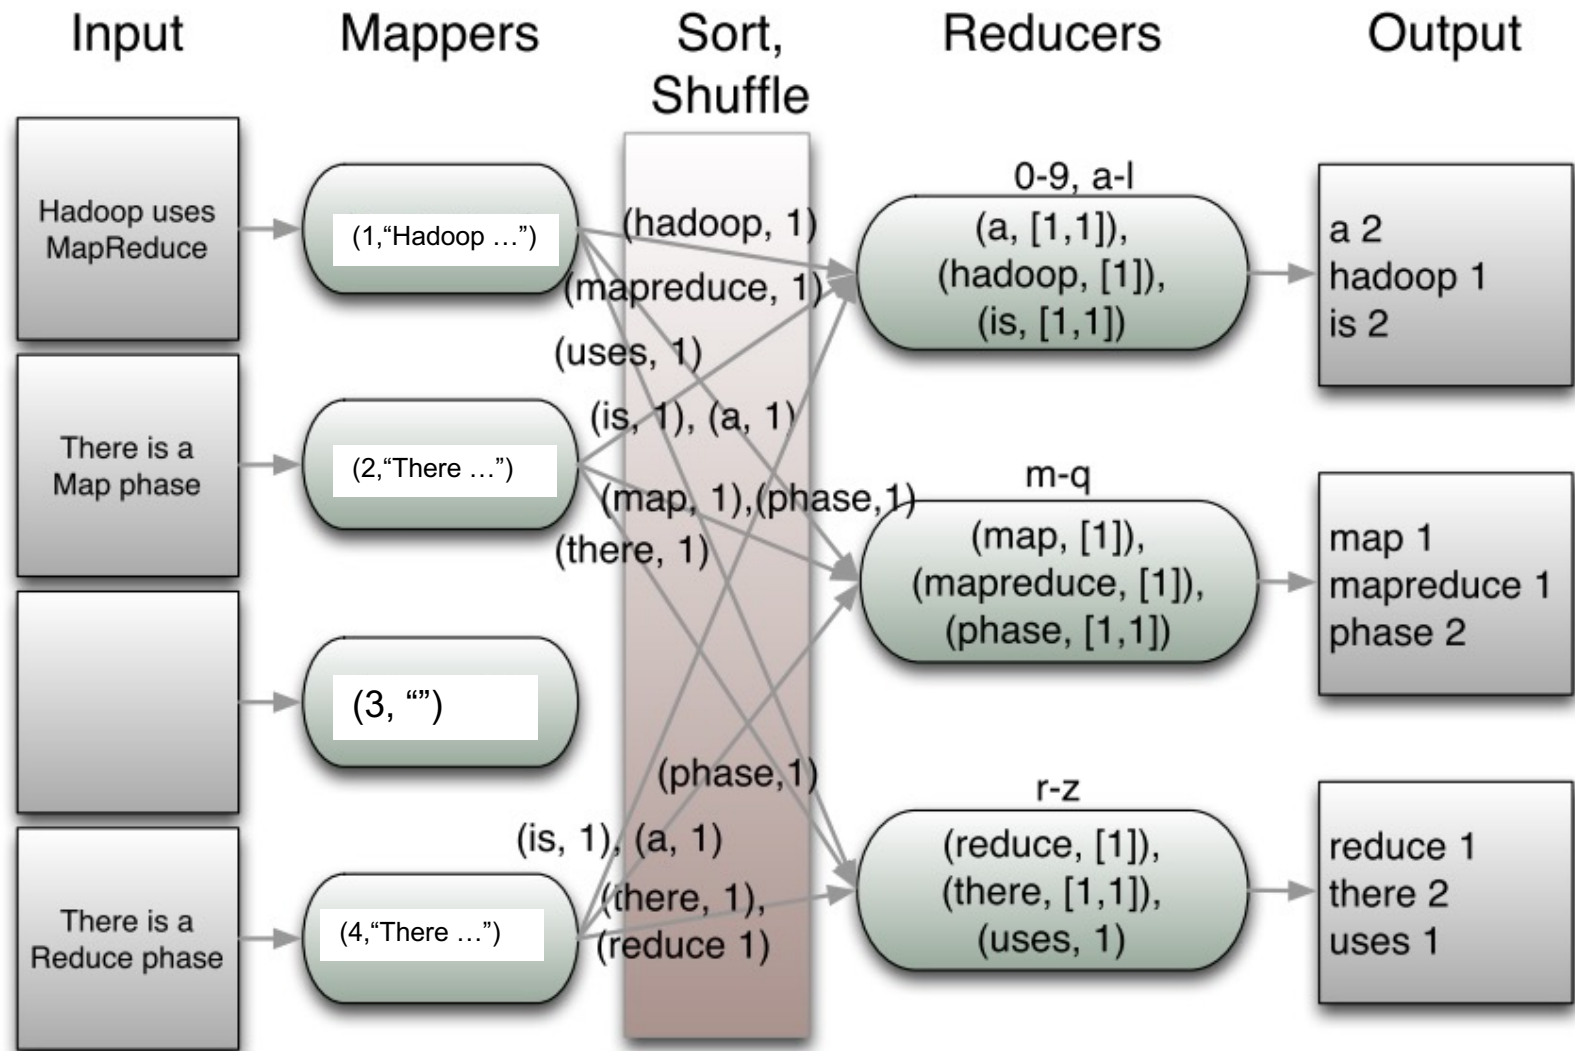
Wordcount Example



Wordcount Example



Wordcount Example



MapReduce Programming Model

For every input line (record), call a map function on the line.

Map:

- Accepts *input* key/value pair
- Emits *intermediate* key/value pairs

For every distinct intermediate key, call a reduce function

Reduce :

- Merges all *intermediate* key/value pairs with the same intermediate key
- Emits *output* key/value pairs (usually one)

Wordcount's Map and Reduce Functions (Pseudo Code)

```
map(key, value):                                     // For 1st line: key = 1, value = "hadoop uses mapreduce"
    // key: line number
    // value: line content
    for each word w in split(value, " ") ← Split returns an array of words
        Emit(w, 1);                                // Split() returns ["hadoop", "uses", "mapreduce"]
    end for                                         // Emit("hadoop",1)
                                                    // Emit("uses", 1)
                                                    // Emit("mapreduce",1)
```

```
reduce(key, values[ ]):                             // For 1st group: key = "a", value = [1,1]
    // key: a word
    // values: a list of 1 for the same word
    sum = 0;
    for each v in values
        sum += v;
    end for
    Emit(key, sum);                                // Emit("a",2)
```

Mean Temperature Example

- From data collected since 2010 until now, find the mean temperature of each month
- Input data has two columns: YYYYMMDDhhmm, Celsius
- Output has two columns and 12 rows each for a month

```
201001010000, 25.0
201001010015, 24.5
201001010030, 24.0
201001010045, 24.0
201001010100, 23.5
201001010115, 23.0
...
201512312345, 27.0
```


Mean's Map and Reduce Functions

```
map(key, value):
```

```
    // key: line number
```

```
    // value: line of datetime and degree
```

```
    (datetime,degree) = split(value, ",") ←
```

```
    month = substring(datetime,4,2)
```

```
    Emit(month, degree);
```

Split to array and then assign
the first and second elements to
datetime and degree

```
    // For the first 2 lines
```

```
    // Emit ("01", "25.0")
```

```
    // ("01", "24.5")
```

```
reduce(key, values[ ]):
```

```
    // For 1st group: key="01" , value = ["25.0", "24.5". ...]
```

```
    // key: a month
```

```
    // values: a list of temperatures in the same month
```

```
    sum = 0, count=0;
```

```
    for each v in values
```

```
        sum += v;
```

```
        count += 1;
```

```
    end for
```

```
    Emit(key, sum/count);
```

Searching Keyword Example

- Example: Check if a specific keyword “*refund*” exists in a large text file

Some people believe that getting a large tax **refund** is not as desirable as more accurate withholding throughout the year, as a large **refund** represents a loan paid back by the government interest-free. Optimally, a return should result in a payment owed of just less than would cause a penalty charge, which is 100% of the prior year's tax (110% for high income individuals), 90% of the current year's tax, or \$1,000 for individuals who have direct withholding and do not pay estimated tax. In order to decrease the amount of the tax **refund** which has to be received by taxpayers, they can turn to one or several of the following methods ...

Searching's Map and Reduce Functions

```
map(key, value):  
    // key: line number  
    // value: line of text  
    count = 0  
    for each word w in split(value, " ")  
        if (w equals "refund") count += 1  
    end for  
    Emit("refund", count)    // 0 means not found
```

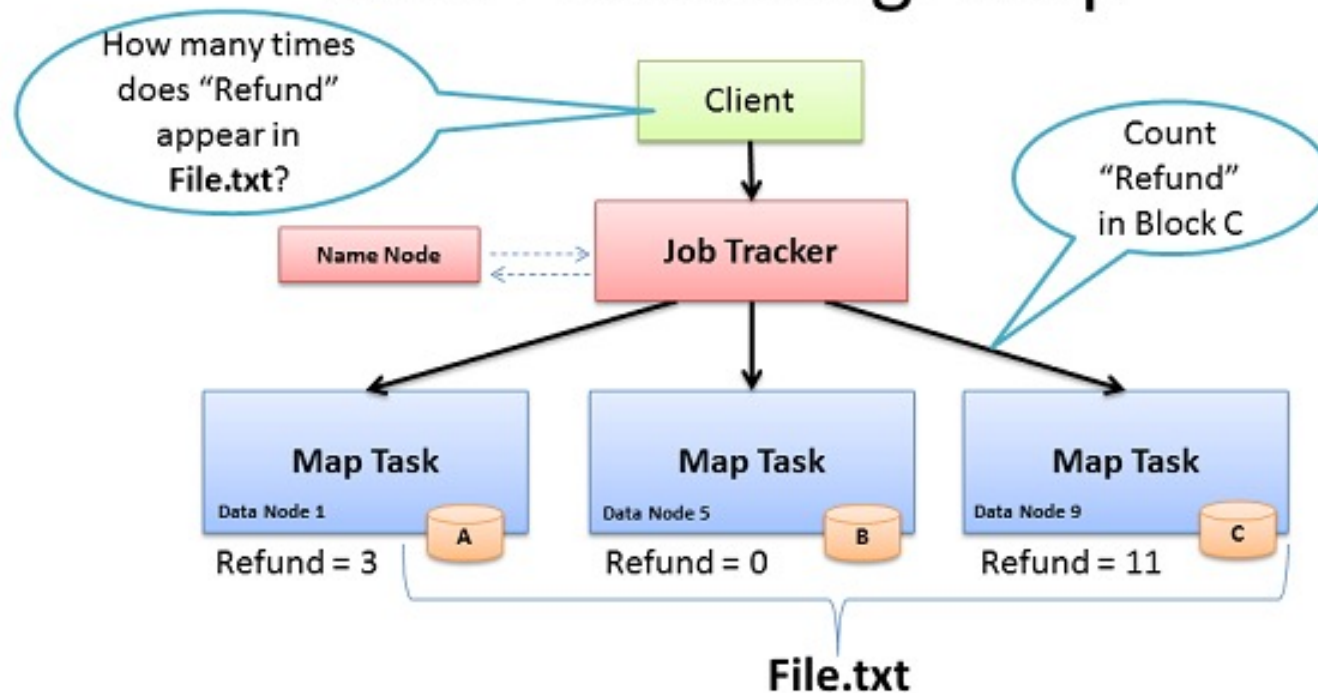
```
reduce(key, values[ ]):  
    // key: word "refund"  
    // values: a list of frequencies for the word "refund" in each line  
    sum = 0;  
    for each v in values  
        sum += v;  
    end for  
    Emit(key, sum); // If output has sum > 0, found
```

Hadoop Job Scheduling

- Input files are divided into blocks (default to HDFS blocks)
- A map task is created for each block
- JobTracker finds the location of block from NameNode
- JobTracker selects a node to run a map task
- Data aware scheduling
 - A map task is scheduled to run at the node where data block is located
 - Otherwise, a map task is scheduled to run at the node close to the data block

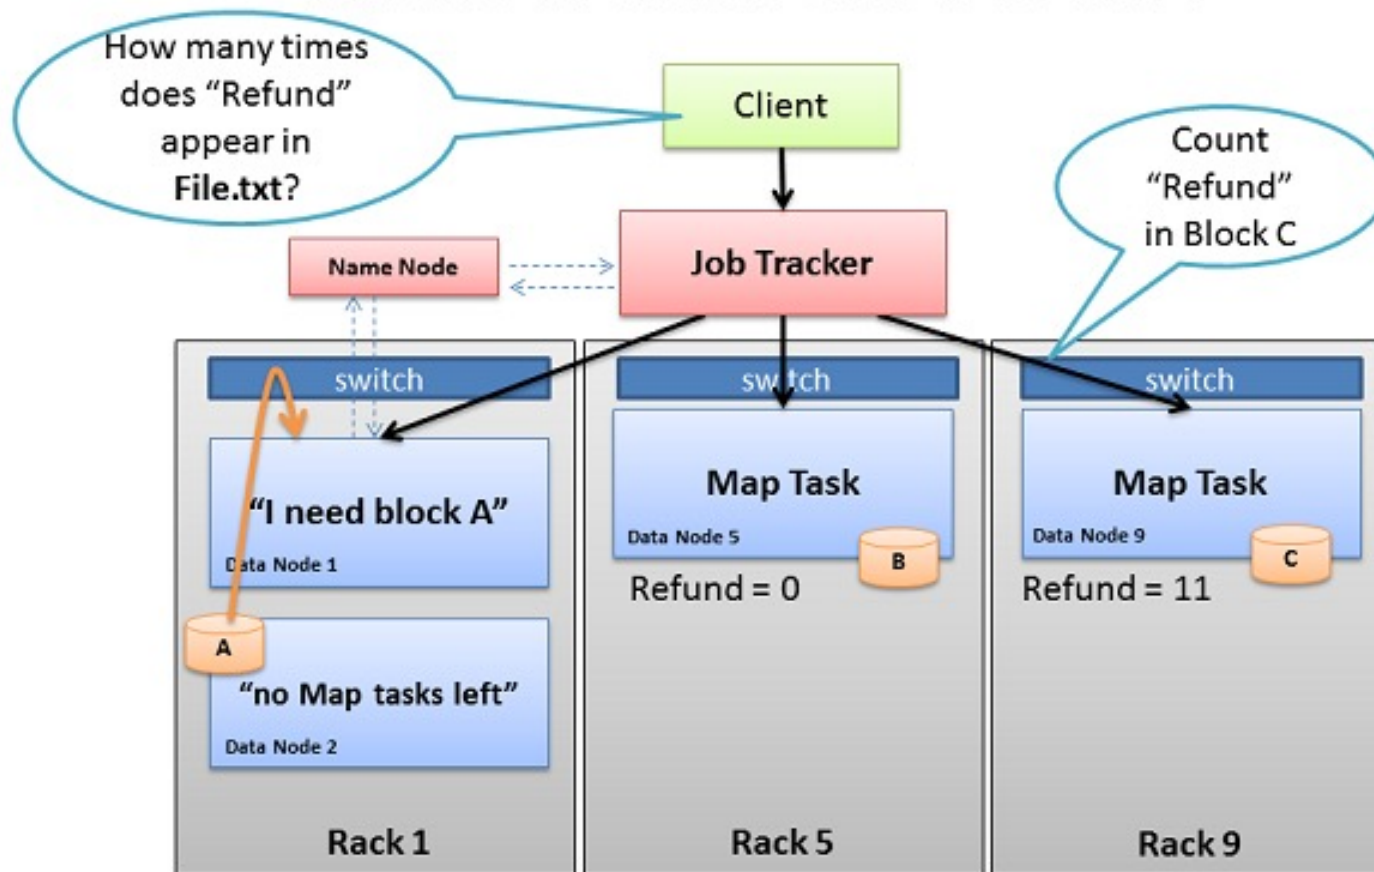
Data Aware Scheduling in Map Tasks

Data Processing: Map



- **Map:** "Run this computation on your local data"
- Job Tracker delivers Java code to Nodes with local data

What if data isn't local?



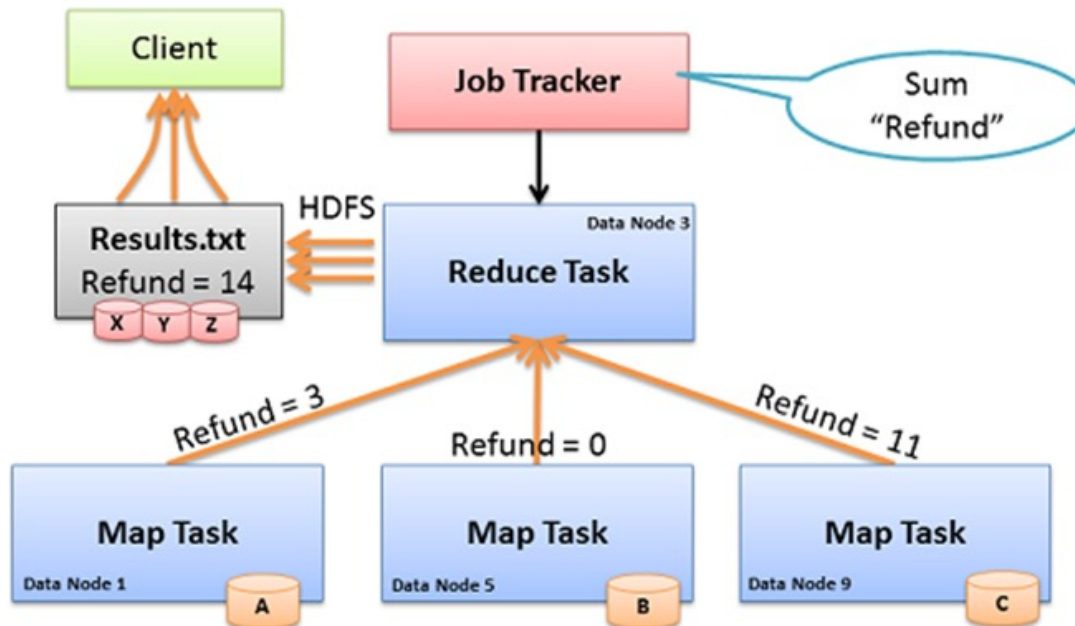
- Job Tracker tries to select Node in same rack as data
- Name Node rack awareness

BRAD HEDLUND .com

Reduce Task Scheduling

- Reduce tasks are placed at random nodes
- The intermediate data from mappers are split to each reducer based on the **hash or range partition** of the intermediate key
- For wordcount example, with 2 reducers, using range partition, words begin with a-m is assigned to reducer1, and n-z is assigned to reducer2
- Each reduce task collects intermediate data of the assigned keys from all mappers

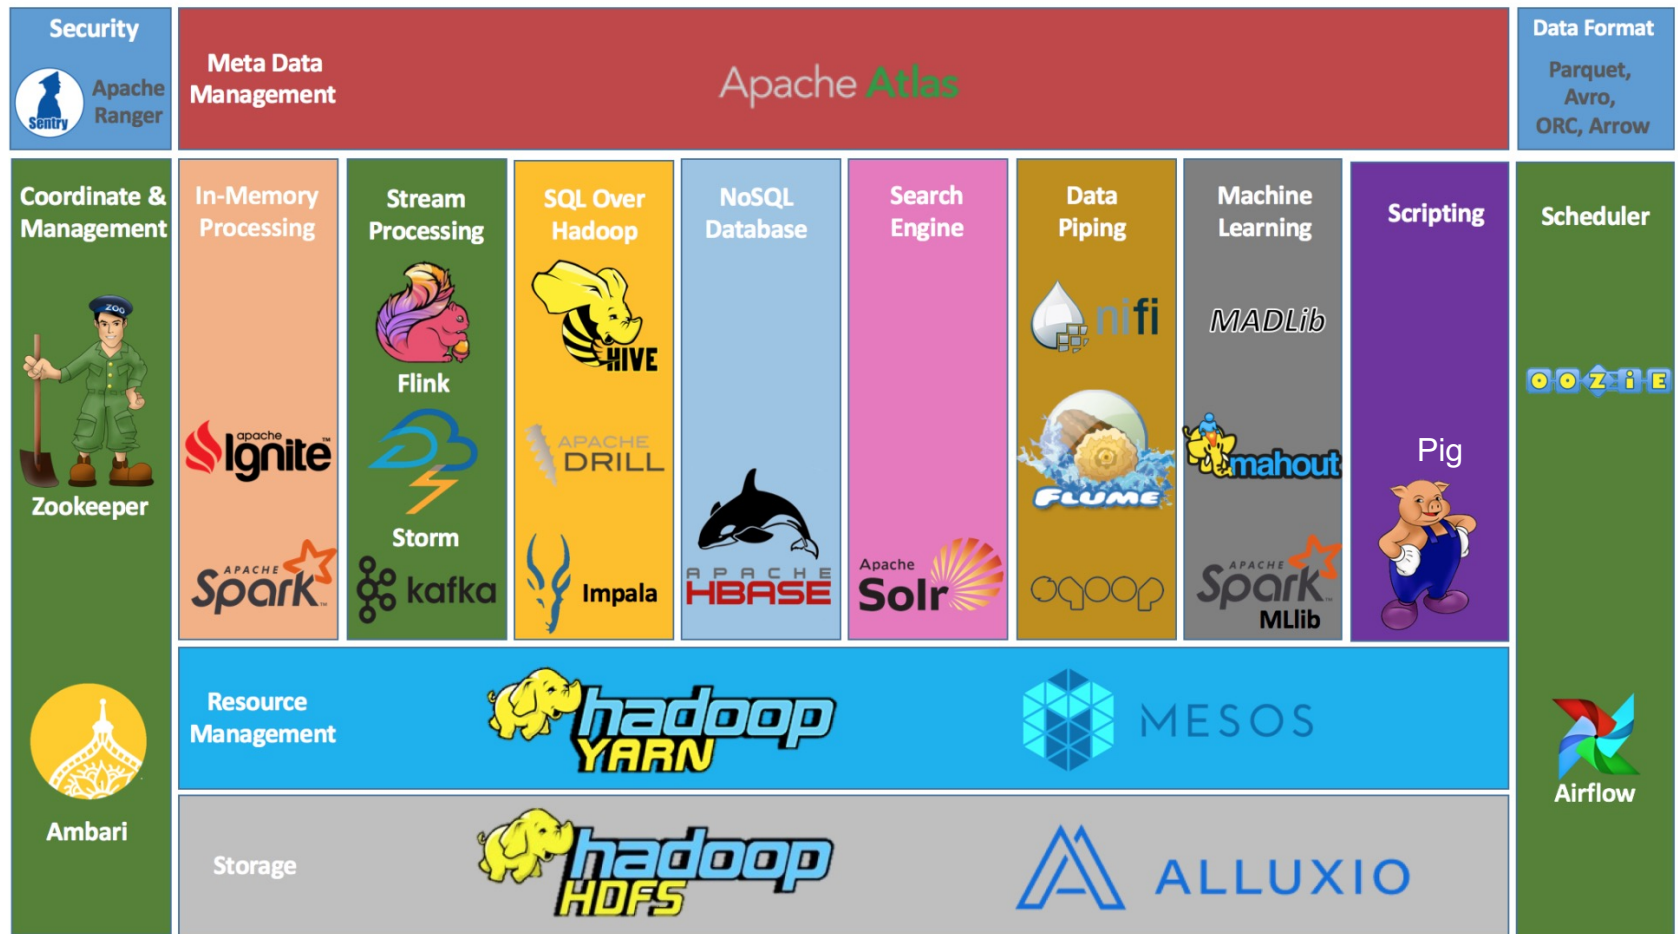
Data Processing: Reduce



- **Reduce:** “Run this computation across Map results”
- Map Tasks send output data to Reducer over the network
- Reduce Task data output written to HDFS

BRAD HEDLUND .com

Hadoop Ecosystem



HBASE

- A NoSQL datastore

- Clone of Big Table (Google)
- Implemented in Java (Clients : Java, C++, Ruby...)
- Data is stored “Column-oriented”
- Distributed over many servers
- Tolerant of machine failure
- Layered over HDFS
- Strong consistency
- It's not a relational database (No joins)
- Sparse data – nulls are stored for free
- Semi-structured or unstructured data
- Data changes through time
- Versioned data
- Scalable – Goal of billions of rows x millions of columns

Table - Example

Row	Timestamp	Animal		Repair
		Type	Size	Cost
Enclosure1	12	Zebra	Medium	1000€
	11	Lion	Big	
Enclosure2	13	Monkey	Small	1500€

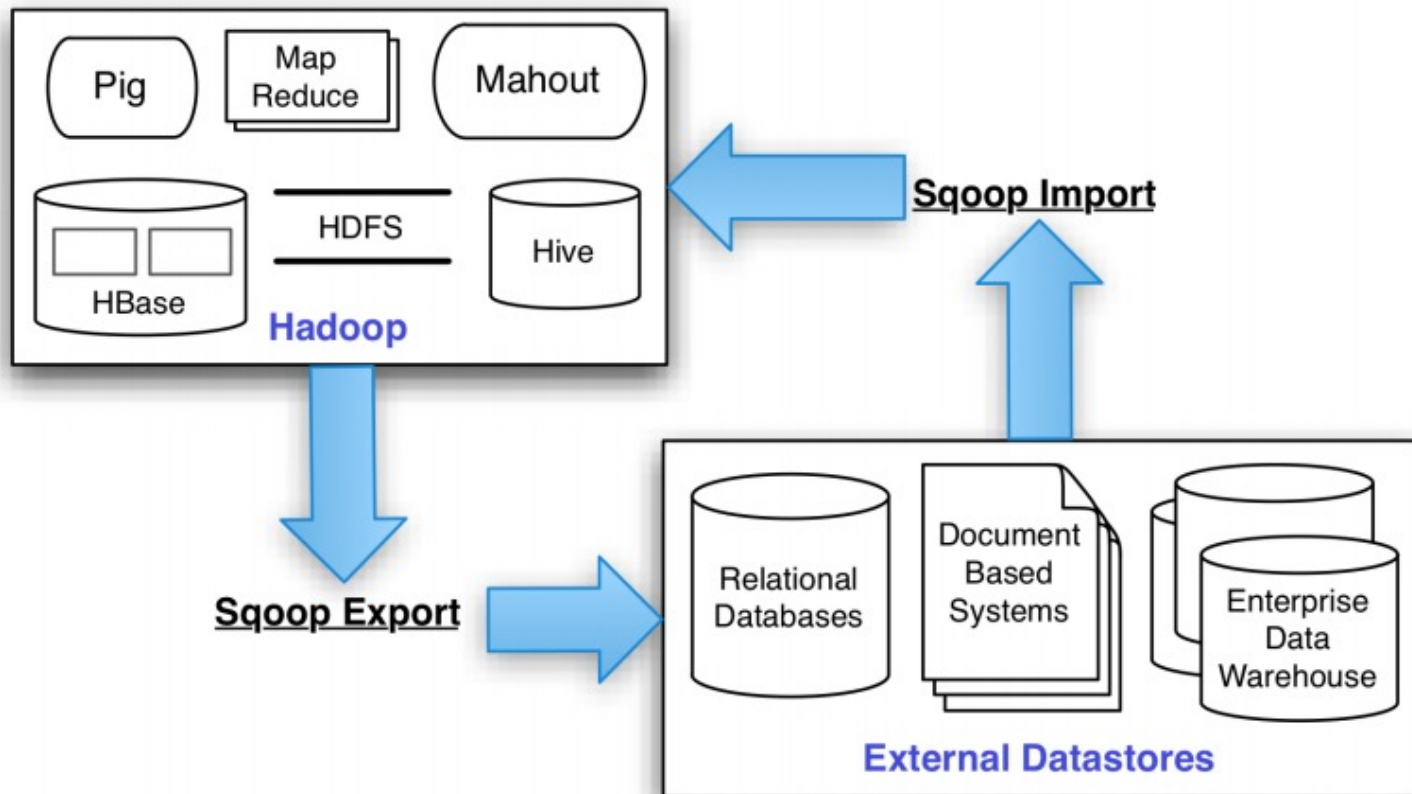
Diagram illustrating the HBase table structure and components:

- Region**: A bracket on the left side of the table indicates the data is distributed across regions.
- Key**: Points to the Row column (Enclosure1, Enclosure2).
- Column**: Points to the Timestamp column (12, 11, 13).
- Family**: Points to the Animal column (Zebra, Lion, Monkey).
- Cell**: Points to the Repair column (1000€, 1500€).

(Table, Row_Key, Family, Column, Timestamp) = Cell (Value)

Sqoop

- A tool for data import/export between Hadoop and RDBMS



Flume

- Move large amount of streaming event (log data) into Hadoop

Access log:

127.0.0.1 - frank [10/Oct/2000:13:55:36 -0700] "GET /apache_pb.gif HTTP/1.0" 200 2326

Error log:

[Sun Mar 7 20:58:27 2004] [info] [client 64.242.88.10] (104)Connection reset by peer: client stopped connection before send body completed

[Sun Mar 7 21:16:17 2004] [error] [client 24.70.56.49] File does not exist:
/home/httpd/twiki/view/Main/WebHome


Vmstat

```
procs -----memory----- ---swap-- -----io----- --system-- -----cpu-----  
r b swpd free buff cache si so bi bo in cs us sy id wa st  
0 0 305416 260688 29160 2356920 2 2 4 1 0 0 6 192 2 0
```

iostat

Linux 2.6.32-100.28.5.el6.x86_64 (dev-db) 07/09/2011


```
avg-cpu:  %user  %nice %system %iowait  %steal   %idle  
          5.68   0.00   0.52   2.03   0.00  91.76
```






- Supersede Hadoop MapReduce
- Can utilize both in-memory and disk storage

Hadoop MapReduce: Data Sharing on Disk



Spark: Speed up processing by using Memory instead of Disks

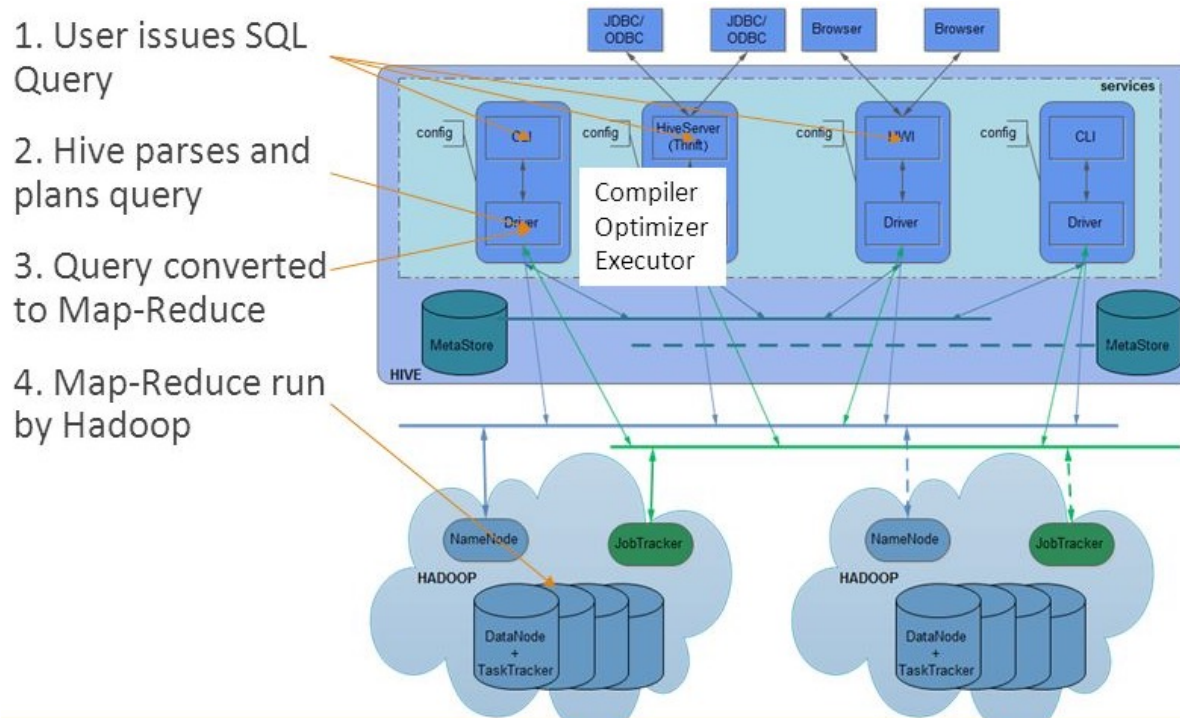


- More operations than map/reduce
- Support Scala, Java, Python, R
- Scale from a single node to a cluster

Hive

- Hive is an SQL-like interface to Hadoop

```
SELECT * FROM purchases WHERE price > 10000 ORDER BY storeid
```



Additional References

- <https://hadoop.apache.org/>
- <http://biforbeginners.blogspot.com/2012/08/hadoop-basics.html>

Q & A

